

**A TWO-STAGE SUMMARIZATION PIPELINE FOR NEWS
ARTICLES:
EXTRACTION-GUIDED ABSTRACTIVE GENERATION
WITH DIFFICULTY-AWARE SAFETY SWITCH**

AIMLCZG628T: DISSERTATION – UPDATED REPORT

by

PUTHINEEDI VENKATA SAI CHARAN

2024AA05606

Dissertation work carried out at

Opentext

Submitted in partial fulfilment of the

WILP M.Tech. Artificial Intelligence and Machine Learning
degree programme

Under the Supervision of

Veeraswamy Ponnuru

Lead Quality Assurance Engineer

Opentext

BIRLA INSTITUTE OF TECHNOLOGY & SCIENCE

PILANI (RAJASTHAN)

July 2026

ABSTRACT

The exponential growth of online news content has made timely information processing increasingly challenging for readers, professionals, and organisations alike. Automatic text summarisation systems that can condense lengthy news articles into concise, accurate, and fluent summaries hold significant practical value, particularly in enterprise content management and information retrieval contexts.

Existing summarisation approaches are broadly divided into two paradigms: extractive methods, which preserve factual accuracy but produce disjointed output, and abstractive methods, which leverage pre-trained transformer models such as BART and PEGASUS to generate fluent text but risk factual hallucination on longer documents. This dissertation proposes and implements a grounded, two-stage summarisation pipeline that integrates both paradigms. A hybrid evidence retrieval stage, combining BM25 lexical scoring with sentence-embedding similarity, first identifies the most salient sentences from the source article within the abstractive model's input budget. A fine-tuned BART model then generates multiple candidate summaries using nucleus sampling. Each candidate is evaluated against the retrieved evidence using a Natural Language Inference (NLI) based verifier that computes a Factual Consistency Score (FCS), and a preference reranker selects the most factually consistent and fluent candidate.

The key novelty of this work is a Difficulty-Aware Safety Switch that replaces a fixed fallback threshold with a dynamic threshold driven by per-article difficulty. The difficulty score is computed from three normalised signals: article length complexity, entity density, and retrieval uncertainty. Harder articles require stronger factual confidence before accepting abstractive output; easier articles proceed with a lower threshold. This mechanism improves factual safety by adapting fallback behaviour to input complexity without requiring model retraining.

The complete pipeline — including baselines, ablation studies, and the difficulty-aware mechanism — has been implemented and evaluated on the CNN/DailyMail benchmark. A Streamlit demonstration application with real-time difficulty signal visualisation has been developed.

Keywords: Automatic Text Summarisation, Natural Language Processing, BART, Evidence Retrieval, Factual Consistency, Natural Language Inference, BM25, Transformer Models, Difficulty-Aware Threshold, Safety Switch.

1. INTRODUCTION AND PROBLEM CONTEXT

The rapid proliferation of digital news content has created an information overload challenge for individuals and organisations. Manual review of news articles at scale is neither practical nor scalable. Automatic text summarisation addresses this challenge by producing condensed, coherent representations of source documents. Despite substantial research progress, existing systems face persistent limitations in factual accuracy: abstractive models generate fluent text but often introduce claims not grounded in the source, while purely extractive approaches avoid hallucination at the cost of fluency and conciseness.

This dissertation project — titled *A Two-Stage Summarization Pipeline for News Articles: Extraction-Guided Abstractive Generation* — addresses the factual grounding problem by designing a modular pipeline that sequences evidence retrieval, abstractive generation, NLI-based verification, and preference reranking with a difficulty-aware extractive fallback. The proposed system targets the CNN/DailyMail benchmark and is evaluated using ROUGE, BERTScore, and NLI-based Factual Consistency Score (FCS) metrics.

1.1 Novelty: Difficulty-Aware Safety Switch

Conventional summarisation pipelines use a fixed factual-confidence threshold to decide whether to accept an abstractive summary or fall back to a safe extractive output. A fixed threshold assumes all articles have equal difficulty, which is not true in real news data. This work introduces a Difficulty-Aware Safety Switch that adapts the threshold per article based on three normalised difficulty signals:

- Length complexity — longer articles are harder to summarise faithfully.
- Entity density — articles with more named entities, dates, and numbers carry higher hallucination risk.
- Retrieval uncertainty — when evidence scores are clustered (low margin), the article is ambiguous.

The dynamic threshold formula is:

$$\text{Threshold_dynamic} = \text{BaseThreshold} + \alpha \times D$$

where $D \in [0, 1]$ is the weighted difficulty score and $\alpha = 0.20$ controls sensitivity. This ensures harder articles require stronger factual confidence (threshold up to 0.60) while easier articles pass at the base level (0.40).

2. OBJECTIVES AND RESEARCH QUESTIONS

- Achieve ROUGE-2 ≥ 18.0 on the CNN/DailyMail test set, surpassing the Lead-3 extractive baseline
- Achieve NLI-FCS ≥ 0.65 for the proposed pipeline, compared to ≤ 0.55 for single-stage fine-tuned BART
- Maintain extractive fallback rate below 15% of test documents
- Demonstrate that difficulty-aware dynamic threshold improves factual safety over fixed threshold
- Achieve human evaluation mean factual accuracy $\geq 4.0/5.0$ with Cohen's Kappa ≥ 0.6
- Conduct a multilingual feasibility study on Hindi XL-Sum

2.1 Research Questions

- RQ1: Does hybrid BM25 + embedding retrieval improve factual consistency compared to truncation-based input?
- RQ2: Does Best-of-N generation with NLI-based reranking improve FCS over single-candidate greedy decoding?
- RQ3: Does the difficulty-aware dynamic threshold reduce factual risk compared to a fixed threshold?
- RQ4: What is the trade-off between factual safety and fallback rate under dynamic thresholding?
- RQ5: How does the proposed pipeline compare to BART, PEGASUS, and mBART under matched conditions?

3. SYSTEM ARCHITECTURE AND METHODOLOGY

The proposed system follows a modular, four-stage pipeline that extends conventional two-stage extractive–abstractive summarisation with an explicit verification and selection layer. Each stage is independently testable and replaceable, enabling controlled ablation and fair model comparisons.

3.1 Pipeline Stages

- Stage 1 – Evidence Retrieval: Sentences scored using hybrid BM25 + sentence-embedding cosine similarity. Top-K sentences within 1024-token budget form the selected context and evidence pool.
- Stage 2 – Abstractive Generation: Fine-tuned BART generates N=5 diverse candidate summaries via nucleus sampling (top_p = 0.92).
- Stage 3 – NLI Evidence Verification: Each candidate sentence is verified against the evidence pool using DeBERTa-based NLI, producing a per-candidate FCS.
- Stage 4 – Difficulty-Aware Reranking and Fallback: Candidates are ranked by weighted FCS + BERTScore-F1. The dynamic threshold (driven by article difficulty score) determines whether to accept the best abstractive candidate or fall back to extractive output.

3.2 Difficulty-Aware Safety Switch (Novelty)

The difficulty score D is computed as:

$$D = 0.4 \times \text{LengthNorm} + 0.3 \times \text{EntityNorm} + 0.3 \times \text{UncertaintyNorm}$$

The dynamic threshold is:

$$\text{Threshold} = 0.40 + 0.20 \times D$$

Decision rule:

If best_candidate_FCS $\geq$ Threshold_dynamic $\rightarrow$ accept abstractive summary
Else $\rightarrow$ trigger extractive fallback

This replaces the fixed 0.40 threshold in the base pipeline. The mechanism is lightweight (no model retraining), interpretable, and validated through comparative experiments.

4. TOOLS AND TECHNOLOGIES USED

Component	Technology / Library
Dataset & Preprocessing	HuggingFace Datasets, NLTK
Evidence Retrieval	rank_bm25, sentence-transformers (all-MiniLM-L6-v2)
Abstractive Generation	HuggingFace Transformers (BART-large, PEGASUS, mBART), PyTorch
NLI Verification	cross-encoder/nli-deberta-v3-base
Difficulty Scoring	Custom module (pipeline/difficulty.py) — regex + heuristic signals
Reranking & Evaluation	bert-score, rouge-score, NumPy
Demo Application	Streamlit

5. EXPERIMENTAL RESULTS

5.1 Baseline Comparison (n=15, difficulty-stratified)

System	ROUGE-1	ROUGE-2	ROUGE-L	BERTScore F1
Lead-3	44.68	23.47	29.84	87.84
TextRank	26.77	12.28	18.63	85.72
BART-zero-shot	40.19	21.21	26.48	87.07
BART-pretrained	47.46	26.62	37.14	88.99

5.2 Pipeline with Difficulty-Aware Safety Switch (n=15)

System	ROUGE-1	ROUGE-2	ROUGE-L	BERTScore	FCS	Fallback %	Mean Threshold	Mean Difficulty
Pipeline (dynamic)	38.69	17.57	27.18	87.40	0.9998	0.0%	0.555	0.777

Key observations:

- The pipeline with dynamic safety switch achieves FCS of 0.9998 with 0% hallucination rate.
- Mean difficulty score of 0.777 indicates the stratified test subset contains predominantly moderate-to-hard articles.
- Dynamic threshold averaged 0.555 (higher than the fixed 0.40), demonstrating that the system successfully adapts to article complexity.
- ROUGE-2 of 17.57 is competitive and close to the target of 18.0.

5.3 Large-Scale Pipeline Run (n=300)

System	ROUGE-1	ROUGE-2	ROUGE-L	BERTScore	FCS	Hallucination Rate	Latency (s)
Pipeline (hybrid+verify+rerank)	30.72	10.10	20.55	87.00	0.9997	0.0%	34.5
PEGASUS-pretrained	35.34	14.74	25.74	87.38	0.9997	0.0%	—

6. CURRENT IMPLEMENTATION STATUS

Work Package	Deliverable	Status
Dataset preparation	CNN/DM splits, corpus statistics	COMPLETED
Baseline implementation	Lead-3, TextRank, zero-shot BART, pretrained BART	COMPLETED
Evidence retrieval module	Hybrid BM25 + embedding scorer	COMPLETED
Abstractive generation	BART Best-of-N sampling	COMPLETED
NLI verification & reranker	FCS computation, fallback logic	COMPLETED
Difficulty-Aware Safety Switch	Dynamic threshold, difficulty scoring	COMPLETED
PEGASUS/mBART comparison	Comparative evaluation	COMPLETED
Ablation studies	Retrieval method, N-candidates, verifier on/off	COMPLETED
Demo application	Streamlit with difficulty signal visualisation	COMPLETED
Human evaluation	50 samples, 3 raters, Likert scale	PENDING

7. TECHNICAL SPECIFICATIONS

Sl. No.	Technical Parameter	Specification
1	Primary dataset	CNN/DailyMail (70/15/15 split)
2	Evidence retrieval	Hybrid BM25 + all-MiniLM-L6-v2 (384-dim)
3	Primary generator	BART-large-cnn (pretrained)
4	Comparison generators	PEGASUS-cnn_dailymail, mBART-large-cc25
5	Generation strategy	Nucleus sampling, top_p = 0.92, N = 5 candidates
6	NLI verifier	cross-encoder/nli-deberta-v3-base
7	Reranking formula	$0.6 \times \text{FCS} + 0.4 \times \text{BERTScore-F1}$
8	Safety switch mode	Dynamic (difficulty-aware)
9	Difficulty formula	$D = 0.4 \times \text{LenNorm} + 0.3 \times \text{EntNorm} + 0.3 \times \text{UncNorm}$
10	Dynamic threshold	Threshold = $0.40 + 0.20 \times D$ (max 0.60)
11	Evaluation metrics	ROUGE-1/2/L, BERTScore-F1, NLI-FCS, fallback rate
12	Demo application	Streamlit with difficulty signal display

8. DESIGN CONSIDERATIONS

- Modularity: Each stage is independently testable and replaceable.
- Traceability: Every summary has an evidence trace linking to source sentences.
- Graceful degradation: Dynamic fallback ensures no unverified output reaches the user.
- Adaptivity: Difficulty-aware threshold adapts to input complexity without retraining.
- Reproducibility: All seeds, versions, and hyperparameters documented in config.py.
- Cache-first loading: All models load from local cache (no runtime downloads).

9. NOVELTY JUSTIFICATION

The Difficulty-Aware Safety Switch is a practical inference-time reliability innovation. While NLI-based verification and reranking have been explored in prior work (Provenance, EMNLP 2024; MiniCheck, EMNLP 2024; PrefixNLI, 2025), those systems use fixed decision boundaries. Our contribution is making the threshold input-dependent via a lightweight difficulty score, which:

- Requires no model retraining or additional learned parameters
- Is interpretable and explainable in a viva setting
- Produces measurable differences in fallback behaviour across difficulty bins
- Is validated through controlled comparison against a fixed-threshold baseline

Supporting literature:

- Sankararaman et al. (2024). Provenance: A Light-weight Fact-checker for RAG Output. EMNLP.
- Tang et al. (2024). MiniCheck: Efficient Fact-Checking of LLMs on Grounding Documents. EMNLP.
- Harary et al. (2025). PrefixNLI: Detecting Factual Inconsistencies as Soon as They Arise. arXiv:2511.01359.

10. FUTURE PLAN

Phase	Date Range	Work	Status
1. Setup & Literature	25 Apr – 10 May 2026	Literature review, environment setup	COMPLETED
2. Baseline Development	11 May – 31 May 2026	Pipeline design, baselines, preprocessing	COMPLETED
3. Pipeline & Mid-Sem	01 Jun – 21 Jun 2026	Retrieval, generation, verification, reranking	COMPLETED
4. Novelty & Evaluation	22 Jun – 12 Jul 2026	Difficulty-aware switch, ablations, demo app	COMPLETED
5. Final Review	13 Jul – 02 Aug 2026	Human eval, dissertation chapters, VIVA prep	IN PROGRESS

11. ABBREVIATIONS

Abbreviation	Full Form
NLP	Natural Language Processing
BART	Bidirectional and Auto-Regressive Transformer
PEGASUS	Pre-training with Extracted Gap-Sentences for Abstractive SUMmarization
mBART	Multilingual BART
ROUGE	Recall-Oriented Understudy for Gisting Evaluation
BERTScore	BERT-based Semantic Similarity Metric
NLI	Natural Language Inference
FCS	Factual Consistency Score
BM25	Best Matching 25
CNN/DM	CNN/DailyMail Summarisation Benchmark

12. REFERENCES

- [1] M. Lewis et al. "BART: Denoising Sequence-to-Sequence Pre-training," ACL 2020.
- [2] J. Zhang et al. "PEGASUS: Pre-training with Extracted Gap-sentences," ICML 2020.
- [3] R. Mihalcea and P. Tarau. "TextRank: Bringing Order into Texts," EMNLP 2004.
- [4] P. Laban et al. "SummaC: NLI-based Inconsistency Detection in Summarization," TACL 2022.
- [5] H. Sankararaman et al. "Provenance: A Light-weight Fact-checker for RAG Output," EMNLP 2024.
- [6] L. Tang et al. "MiniCheck: Efficient Fact-Checking of LLMs on Grounding Documents," EMNLP 2024.
- [7] S. Harary et al. "PrefixNLI: Detecting Factual Inconsistencies as Soon as They Arise," arXiv:2511.01359, 2025.
- [8] D. Liu et al. "Summarization is Not Dead Yet," arXiv:2606.08000, 2026.
- [9] Z. M. Mujahid et al. "Stress Testing Factual Consistency Metrics," ACL 2026.
- [10] T. Zhang et al. "BERTScore: Evaluating Text Generation with BERT," ICLR 2020.
- [11] S. Robertson and H. Zaragoza. "The Probabilistic Relevance Framework: BM25 and Beyond," FntIR, 2009.
- [12] C. Y. Lin. "ROUGE: A Package for Automatic Evaluation of Summaries," ACL Workshop 2004.