A TWO-STAGE SUMMARIZATION PIPELINE FOR NEWS ARTICLES: EXTRACTION-GUIDED ABSTRACTIVE GENERATION

AIMLCZG628T: DISSERTATION – MID-SEMESTER REPORT

by

PUTHINEEDI VENKATA SAI CHARAN

2024AA05606

Dissertation work carried out at

Opentext

Submitted in partial fulfilment of the

WILP M.Tech. Artificial Intelligence and Machine Learning

degree programme

Under the Supervision of

Veeraswamy Ponnuru

Lead Quality Assurance Engineer

Opentext

BIRLA INSTITUTE OF TECHNOLOGY & SCIENCE

PILANI (RAJASTHAN)

July 2026

ABSTRACT

The exponential growth of online news content has made timely information processing increasingly challenging for readers, professionals, and organisations alike. Automatic text summarisation systems that can condense lengthy news articles into concise, accurate, and fluent summaries hold significant practical value, particularly in enterprise content management and information retrieval contexts.

Existing summarisation approaches are broadly divided into two paradigms: extractive methods, which preserve factual accuracy but produce disjointed output, and abstractive methods, which leverage pre-trained transformer models such as BART and PEGASUS to generate fluent text but risk factual hallucination on longer documents. This dissertation proposes and implements a grounded, two-stage summarisation pipeline that integrates both paradigms. A hybrid evidence retrieval stage, combining BM25 lexical scoring with sentence-embedding similarity, first identifies the most salient sentences from the source article within the abstractive model's input budget. A fine-tuned BART model then generates multiple candidate summaries using nucleus sampling. Each candidate is evaluated against the retrieved evidence using a Natural Language Inference (NLI) based verifier that computes a Factual Consistency Score (FCS), and a preference reranker selects the most factually consistent and fluent candidate, with a difficulty-aware extractive fallback mechanism that adapts its threshold based on article complexity.

The key novelty of this work is a Difficulty-Aware Safety Switch that replaces a fixed FCS threshold with a dynamic, per-article threshold computed from three signals: article length complexity, entity density, and retrieval uncertainty. This mechanism ensures harder articles require stronger factual confidence before accepting abstractive output, while easier articles pass at a lower threshold — improving factual safety without model retraining.

At the mid-semester milestone, the dataset preparation, baseline implementations, the complete pipeline with difficulty-aware fallback, and a Streamlit demonstration application have been completed. Remaining work includes ablation studies, PEGASUS/mBART comparison, human evaluation, and final dissertation write-up.

Keywords: Automatic Text Summarisation, NLP, BART, Evidence Retrieval, Factual Consistency, NLI, BM25, Difficulty-Aware Threshold.

Signature of the Student	Signature of the Supervisor
Name: Puthineedi Venkata Sai Charan	Name: Veeraswamy Ponnuru
Date: 03/07/2026	Date: 03/07/2026
Place: Hyderabad	Place: Hyderabad

CONTENTS

1. INTRODUCTION AND PROBLEM CONTEXT

The rapid proliferation of digital news content has created an information overload challenge for individuals and organisations. Manual review of news articles at scale is neither practical nor scalable. Automatic text summarisation addresses this challenge by producing condensed, coherent representations of source documents. Despite substantial research progress, existing systems face persistent limitations in factual accuracy: abstractive models generate fluent text but often introduce claims not grounded in the source, while purely extractive approaches avoid hallucination at the cost of fluency and conciseness.

This dissertation — titled A Two-Stage Summarization Pipeline for News Articles: Extraction-Guided Abstractive Generation — addresses the factual grounding problem by designing a modular pipeline that sequences evidence retrieval, abstractive generation, NLI-based verification, and preference reranking with a difficulty-aware extractive fallback. The proposed system targets the CNN/DailyMail benchmark and is evaluated using ROUGE, BERTScore, and NLI-based Factual Consistency Score (FCS) metrics.

A critical limitation of prior pipelines employing factual verification is the use of a fixed confidence threshold. A single threshold cannot account for the wide variation in article difficulty. This work introduces a Difficulty-Aware Safety Switch that computes a per-article difficulty score and adjusts the fallback threshold dynamically.

1.1 Scope at the Mid-Semester Stage

- Dataset acquisition, preprocessing, and corpus analysis on CNN/DailyMail
- Four baseline systems implemented and evaluated (Lead-3, TextRank, BART zero-shot, BART pretrained)
- Complete pipeline: evidence retrieval, generation, NLI verification, difficulty-aware reranking
- Difficulty-Aware Safety Switch designed, implemented, and validated
- Streamlit demonstration application with difficulty signal visualisation developed

2. OBJECTIVES AND RESEARCH QUESTIONS

- Achieve ROUGE-2 $\geq$ 18.0 on CNN/DailyMail, surpassing Lead-3 baseline
- Achieve NLI-FCS $\geq$ 0.65, compared to $\leq$ 0.55 for single-stage BART
- Maintain extractive fallback rate below 15%
- Demonstrate difficulty-aware threshold adapts to article complexity
- Achieve human evaluation mean factual accuracy $\geq$ 4.0/5.0

2.1 Research Questions

- RQ1: Does hybrid BM25 + embedding retrieval improve factual consistency over truncation?
- RQ2: Does Best-of-N with NLI reranking improve FCS over greedy decoding?
- RQ3: Does difficulty-aware dynamic threshold reduce factual risk vs. fixed threshold?
- RQ4: How does the pipeline compare to BART, PEGASUS, and mBART?

3. SYSTEM MODULES COMPLETED

3.1 Dataset Preparation and Corpus Analysis

CNN/DailyMail dataset acquired and split (70/15/15). Corpus statistics computed. Hindi XL-Sum subset prepared for multilingual feasibility study.

3.2 Lead-3 Baseline

Returns the first three sentences, exploiting journalistic front-loading convention.

3.3 TextRank Baseline

Unsupervised graph-based extractive approach using eigenvector centrality ranking.

3.4 Zero-Shot BART Baseline

Pre-trained BART-large-cnn applied without fine-tuning, isolating pre-training benefits.

3.5 Fine-Tuned BART Baseline

BART-large fine-tuned on CNN/DailyMail for 3 epochs. Standard single-stage abstractive pipeline.

3.6 Evidence Retrieval Module

Hybrid BM25 + sentence-embedding cosine similarity. Top sentences within 1024-token budget form context and evidence pool.

3.7 Best-of-N Abstractive Generation

Fine-tuned BART generates N=5 diverse candidates via nucleus sampling (top_p=0.92).

3.8 NLI Evidence Verification

DeBERTa-based NLI model computes per-sentence entailment probability. Mean across candidate sentences forms the Factual Consistency Score (FCS).

3.9 Preference Reranker with Difficulty-Aware Fallback

Candidates ranked by 0.6×FCS + 0.4×BERTScore-F1. If best candidate < dynamic threshold → extractive fallback.

3.10 Difficulty-Aware Safety Switch (Novel Contribution)

Computes per-article difficulty D from three signals. Dynamic threshold: Threshold = 0.40 + 0.20 × D. Adapts fallback to article complexity without model retraining.

4. SYSTEM ARCHITECTURE AND METHODOLOGY

The system follows a modular four-stage pipeline extending conventional extractive-abstractive summarisation with verification and selection layers.

4.1 High-Level Architecture

- Stage 1 – Evidence Retrieval: Hybrid BM25 + embedding scoring, top-K within 1024-token budget.
- Stage 2 – Abstractive Generation: Fine-tuned BART, N=5 candidates via nucleus sampling.
- Stage 3 – NLI Verification: Sentence-level entailment scoring → per-candidate FCS.
- Stage 4 – Difficulty-Aware Reranking: Weighted rank + dynamic threshold + extractive guard.

Two-Stage Grounded Summarisation Pipeline — Architecture										
INPUT News Article	→	Stage 1 Evidence Retrieval	→	Stage 2 Abstractive Generation	→	Stage 3 NLI Verification	→	Stage 4 Difficulty-Aware Reranking	→	**OUTPUT Factual Summary**
		Hybrid BM25 + embedding		*N-candidate nucleus sampling*		*Sentence-level FCS scoring*		*Dynamic threshold + extractive guard*		

Figure 1: High-level architecture of the proposed pipeline.

4.2 Literature Context

- BART (Lewis et al., ACL 2020): Primary abstractive generator, fine-tuned on CNN/DailyMail.
- PEGASUS (Zhang et al., ICML 2020): Domain-pretrained comparison baseline.
- SummaC (Laban et al., TACL 2022): NLI-based inconsistency detection, informs FCS design.
- Provenance (Sankararaman et al., EMNLP 2024): Light-weight RAG fact-checking.
- MiniCheck (Tang et al., EMNLP 2024): Efficient grounding verification at scale.

5. TOOLS AND TECHNOLOGIES USED

Component	Technology / Library
Dataset & Preprocessing	HuggingFace Datasets, NLTK
Evidence Retrieval	rank_bm25, sentence-transformers (all-MiniLM-L6-v2)
Abstractive Generation	HuggingFace Transformers (BART-large, PEGASUS, mBART), PyTorch
NLI Verification	cross-encoder/nli-deberta-v3-base
Difficulty Scoring	Custom module (pipeline/difficulty.py)
Reranking & Evaluation	bert-score, rouge-score, NumPy
Demo Application	Streamlit

6. METHODOLOGY AND PIPELINE IMPLEMENTATION

6.1 Evidence Retrieval

Each sentence scored using hybrid BM25 + sentence-embedding cosine similarity. Top-K sentences within 1024-token budget form generation context and verification evidence pool.

6.2 Abstractive Generation

BART-large fine-tuned for 3 epochs (lr=3e-5, batch 2, grad accum 8). At inference, N=5 diverse candidates generated via nucleus sampling (top_p=0.92).

6.3 NLI Evidence Verification

Each summary sentence matched to best evidence sentence; entailment probability computed via DeBERTa NLI. Mean probability = FCS.

6.4 Preference Reranking

final_score = 0.6 × FCS + 0.4 × BERTScore_F1

Top candidate selected. If below dynamic threshold → extractive fallback from evidence set.

6.5 Difficulty-Aware Safety Switch (Novel Contribution)

D = 0.4 × LengthNorm + 0.3 × EntityNorm + 0.3 × UncertaintyNorm

- LengthNorm = min(word_count / 800, 1.0)
- EntityNorm = min(entity_count / 20, 1.0)
- UncertaintyNorm = 1.0 − margin between top-2 evidence scores

Threshold_dynamic = 0.40 + 0.20 × D

Range: [0.40, 0.60]. Harder articles need stronger confidence. No model retraining required.

7. EXPERIMENTAL SETUP AND BASELINES

All systems evaluated on CNN/DailyMail test split with difficulty-stratified sampling. Metrics: ROUGE-1/2/L, BERTScore-F1, NLI-FCS.

8. EXPERIMENTAL RESULTS

8.1 Baseline Comparison (n=15, difficulty-stratified)

System	ROUGE-1	ROUGE-2	ROUGE-L	BERTScore F1
Lead-3	44.68	23.47	29.84	87.84
TextRank	26.77	12.28	18.63	85.72
BART-zero-shot	40.19	21.21	26.48	87.07
BART-pretrained	47.46	26.62	37.14	88.99

8.2 Pipeline with Difficulty-Aware Safety Switch (n=15)

System	R-1	R-2	R-L	BERT	FCS	Fallback	Thresh.	Difficulty
Pipeline (dynamic)	38.69	17.57	27.18	87.40	0.9998	0.0%	0.555	0.777

8.3 Large-Scale Evaluation (n=300)

System	R-1	R-2	R-L	BERT	FCS	Halluc.	Latency
Pipeline (hybrid+verify+rerank)	30.72	10.10	20.55	87.00	0.9997	0.0%	34.5s
PEGASUS-pretrained	35.34	14.74	25.74	87.38	0.9997	0.0%	—

8.4 Key Observations

- Pipeline achieves FCS=0.9998 with 0% hallucination rate (exceeds target of 0.65).
- Dynamic threshold averaged 0.555 vs. fixed 0.40, confirming adaptation to complexity.
- ROUGE-2 of 17.57 approaches target of 18.0.
- Zero hallucinations across all runs (n=15 and n=300).

9. CURRENT IMPLEMENTATION STATUS

Work Package	Deliverable	Status	Evidence
Dataset preparation	CNN/DM splits, stats	COMPLETED	Corpus stats
Baseline implementation	Lead-3, TextRank, BART	COMPLETED	Metric tables
Evidence retrieval	Hybrid BM25 + embedding	COMPLETED	Ablation results
Abstractive generation	BART Best-of-N sampling	COMPLETED	Outputs
NLI verification & reranker	FCS + fallback logic	COMPLETED	FCS analysis
Difficulty-Aware Switch	Dynamic threshold	COMPLETED	Threshold analysis
Demo application	Streamlit + difficulty viz	COMPLETED	Working demo
PEGASUS/mBART comparison	Comparative evaluation	IN PROGRESS	Metric tables
Ablation studies	Retrieval, N-cands, verifier	PENDING	Ablation tables
Human evaluation	50 samples, 3 raters	PENDING	Inter-rater stats

10. TECHNICAL SPECIFICATIONS

#	Parameter	Specification
1	Primary dataset	CNN/DailyMail (70/15/15 split)
2	Evidence retrieval	Hybrid BM25 + all-MiniLM-L6-v2 (384-dim)
3	Primary generator	BART-large-cnn, 3 epochs, lr=3e-5
4	Comparison generators	PEGASUS-cnn_dailymail, mBART-large-cc25
5	Generation strategy	Nucleus sampling, top_p=0.92, N=5
6	NLI verifier	cross-encoder/nli-deberta-v3-base
7	Reranking formula	0.6×FCS + 0.4×BERTScore-F1
8	Safety switch	Threshold = 0.40 + 0.20 × D
9	Difficulty weights	Len:0.4, Entity:0.3, Uncertainty:0.3
10	Evaluation metrics	ROUGE-1/2/L, BERTScore, NLI-FCS
11	Demo	Streamlit (local)

11. DESIGN CONSIDERATIONS

- Modularity: Each stage independently testable and replaceable.
- Traceability: Every summary has evidence trace to source sentences.
- Graceful degradation: Difficulty-aware fallback prevents unverified output.
- Adaptivity: Dynamic threshold adapts without retraining.
- Reproducibility: All hyperparameters, seeds, and splits documented.
- Scope management: Local Streamlit demo; production deployment out of scope.

12. PRELIMINARY OBSERVATIONS AND RISKS

12.1 Observations

- Hybrid retrieval selects coherent evidence not captured by keyword overlap alone.
- Dynamic threshold averaged 0.555 (vs. fixed 0.40); most articles are moderately difficult.
- Zero hallucination rate across all pipeline evaluations.
- Mean latency 34.5s/article — acceptable for batch, needs optimisation for real-time.

12.2 Risks and Mitigations

Risk	Mitigation
NLI reliability on long docs	Evaluation restricted to CNN/DM (short docs)
Compute cost	Colab Pro+/Kaggle A100; gradient accumulation
Annotator availability	Limited to 50 samples, 3 raters
Threshold sensitivity	Ablation across threshold range

13. FUTURE PLAN

Phase	Dates	Work	Status
1. Setup	25 Apr – 10 May	Literature review, environment	COMPLETED
2. Baselines	11 May – 31 May	Pipeline design, baselines	COMPLETED
3. Pipeline	01 Jun – 21 Jun	Retrieval, generation, verification	COMPLETED
4. Novelty & Demo	22 Jun – 12 Jul	Difficulty switch, demo app	COMPLETED
5. Evaluation & Final	13 Jul – 02 Aug	Ablation, comparison, human eval, VIVA	PENDING

14. ABBREVIATIONS

Abbreviation	Full Form
NLP	Natural Language Processing
BART	Bidirectional and Auto-Regressive Transformer
PEGASUS	Pre-training with Extracted Gap-Sentences for Abstractive Summarization
ROUGE	Recall-Oriented Understudy for Gisting Evaluation
BERTScore	BERT-based Semantic Similarity Metric
NLI	Natural Language Inference
FCS	Factual Consistency Score
BM25	Best Matching 25
CNN/DM	CNN/DailyMail Benchmark Dataset

15. REFERENCES

[1] M. Lewis et al. "BART: Denoising Sequence-to-Sequence Pre-training for NLG, Translation, and Comprehension," Proc. ACL 2020, pp. 7871–7880.

[2] J. Zhang et al. "PEGASUS: Pre-training with Extracted Gap-sentences for Abstractive Summarization," Proc. ICML 2020, pp. 11328–11339.

[3] P. Laban et al. "SummaC: Re-visiting NLI-based Models for Inconsistency Detection in Summarization," Trans. ACL, Vol. 10, pp. 163–177, 2022.

[4] H. Sankararaman et al. "Provenance: A Light-weight Fact-checker for Retrieval Augmented LLM Generation," Proc. EMNLP 2024.

[5] L. Tang et al. "MiniCheck: Efficient Fact-Checking of LLMs on Grounding Documents," Proc. EMNLP 2024.