

BITS Pilani — Work Integrated Learning Programmes

Department of Computer Science and Information Systems

M.Tech Dissertation

A Two-Stage Summarisation Pipeline for News Articles

Extraction-Guided Abstractive Generation with Evidence Retrieval,
NLI-based Factual Verification, and a Difficulty-Aware Safety Switch

Submitted by

Puthineedi Venkata Sai Charan

ID: 2024AA05606

Under the Supervision of

Veeraswamy Ponnuru
Lead QA Engineer, Opentext

Course: AIMLCZG628T — Dissertation

Programme: M.Tech (Artificial Intelligence and Machine Learning)

BITS Pilani — WILP Division

August 2025

Declaration

I, **Puthineedi Venkata Sai Charan** (ID: 2024AA05606), hereby declare that the dissertation entitled "*A Two-Stage Summarisation Pipeline for News Articles: Extraction-Guided Abstractive Generation with Evidence Retrieval, NLI-based Factual Verification, and a Difficulty-Aware Safety Switch*" submitted to BITS Pilani, Work Integrated Learning Programmes Division, in partial fulfillment of the requirements for the award of the degree of Master of Technology in Artificial Intelligence and Machine Learning, is a record of original work carried out by me under the supervision of **Veeraswamy Ponnuru**, Lead QA Engineer, Opentext.

I further declare that this dissertation has not been submitted, either in part or in full, for the award of any other degree or diploma of any university or institution.

Place: _____

Date: _____

Signature of Student

Puthineedi Venkata Sai Charan

2024AA05606

Acknowledgements

I would like to express my sincere gratitude to my dissertation supervisor, **Veeraswamy Ponnuru**, Lead QA Engineer, Opentext, for his invaluable guidance, constructive feedback, and continuous encouragement throughout the course of this research. His practical insights into industrial NLP applications greatly shaped the direction of this work.

I am deeply grateful to the faculty and staff of BITS Pilani, Work Integrated Learning Programmes Division, for their academic support and for providing a stimulating research environment. The rigorous curriculum of the M.Tech AIML programme provided me with the theoretical foundations necessary to undertake this research.

I wish to acknowledge the open-source community, in particular the teams behind Hugging Face Transformers, the CNN/DailyMail benchmark dataset maintainers, and the developers of BART, PEGASUS, DeBERTa, and Sentence-Transformers, whose publicly available pretrained models and tools made this research feasible within the resource constraints of a work-integrated programme.

The experiments in this dissertation were conducted on Google Colab with free-tier T4 GPU resources. I acknowledge Google LLC for providing free-tier cloud computation that enabled the training and evaluation of large transformer models within the constraints of this programme.

Finally, I am grateful to my family and colleagues for their patience, moral support, and understanding during the demanding period of this research.

Puthineedi Venkata Sai Charan

August 2025

Abstract

Automatic news summarisation requires generating concise, fluent, and factually consistent summaries of lengthy news articles. While large pre-trained sequence-to-sequence models such as BART and PEGASUS achieve strong ROUGE scores on benchmarks such as CNN/DailyMail, they are susceptible to hallucination — generating plausible-sounding claims that are not supported by the source text. Existing single-stage abstractive pipelines provide no mechanism to verify factual consistency or to fall back to a safer output when generation confidence is low.

This dissertation proposes and evaluates a four-stage extraction-guided abstractive summarisation pipeline that directly addresses these limitations. The pipeline comprises: (1) a hybrid BM25 and sentence-embedding evidence retriever that selects the most salient sentences within a token budget; (2) a Best-of-N nucleus sampling generator based on fine-tuned BART-large that produces multiple candidate summaries; (3) a DeBERTa-based Natural Language Inference (NLI) verifier that assigns a Factual Consistency Score (FCS) to each candidate; and (4) a preference reranker that selects the best candidate by a weighted combination of FCS and BERTScore-F1. A novel Difficulty-Aware Safety Switch dynamically adjusts the fallback threshold based on estimated article difficulty, reducing factual risk on complex inputs while preserving fluency on simpler ones.

The system is evaluated on the CNN/DailyMail 3.0.0 benchmark using ROUGE-1, ROUGE-2, ROUGE-L, BERTScore-F1, NLI-FCS, and fallback rate. Ablation studies systematically isolate the contribution of retrieval strategy (TextRank, BM25-only, embedding-only, hybrid), candidate count (Best-of-1 through Best-of-8), and the NLI verifier. The hybrid retrieval arm achieves ROUGE-2 of 24.18 on a 15-sample development set, outperforming BM25-only (12.78) and embedding-only (19.08) retrievers. Best-of-8 sampling reaches ROUGE-2 of 21.36 versus 17.81 for single-candidate generation. The full pipeline achieves an NLI-FCS of 99.97% on a 300-sample evaluation, with zero hallucinations detected and a zero percent extractive fallback rate, indicating high-confidence abstractive generation across all test documents. A multilingual feasibility study demonstrates the viability of extending the pipeline to Hindi using mBART fine-tuned on the XL-Sum dataset.

The principal contributions of this dissertation are: a modular, reproducible four-stage pipeline that combines retrieval, generation, verification, and reranking in a single framework; a novel Difficulty-Aware Safety Switch for inference-time factual risk management; systematic ablation evidence justifying every design choice; and a demonstration that a carefully engineered pipeline can achieve near-perfect factual consistency on standard news summarisation benchmarks without requiring reinforcement learning from human feedback.

Keywords: Abstractive Summarisation, BART, Evidence Retrieval, Factual Consistency, NLI Verification, Difficulty-Aware Safety Switch, CNN/DailyMail, Best-of-N, BERTScore, Hybrid Retrieval

Table of Contents

Declaration	i
Acknowledgements	ii
Abstract	iii
Table of Contents	iv
List of Figures	v
List of Tables	vi
Chapter 1 Introduction	1
1.1 Background and Motivation	1
1.2 Research Problem	2
1.3 Objectives	2
1.4 Research Questions	3
1.5 Scope and Constraints	3
1.6 Contributions	3
1.7 Report Organisation	4
Chapter 2 Literature Review	5
2.1 Extractive Summarisation	5
2.2 Abstractive Summarisation	5
2.3 Pre-trained Transformer Models	6
2.4 Factuality Evaluation	6
2.5 Retrieval-Augmented Generation	7
2.6 Research Gap and Contribution	7
Chapter 3 Problem Statement and Objectives	9
3.1 Problem Statement	9
3.2 Research Objectives	9
3.3 Research Questions	10
3.4 Scope	10
Chapter 4 System Design and Architecture	11
4.1 High-Level Architecture	11
4.2 Dataset: CNN/DailyMail	11
4.3 Stage 1 — Evidence Retriever	12
4.4 Stage 2 — Abstractive Generator	13
4.5 Stage 3 — NLI Verifier	14
4.6 Stage 4 — Preference Reranker	15
4.7 Difficulty-Aware Safety Switch	15

4.8 Technology Stack	16
Chapter 5 Implementation	17
5.1 Repository Structure	17
5.2 Data Pipeline	17
5.3 Model Training	18
5.4 Evaluation Harness	19
5.5 Streamlit Demo Application	19
5.6 Multilingual Extension	19
Chapter 6 Experimental Evaluation	20
6.1 Evaluation Metrics	20
6.2 Baseline Results	21
6.3 Ablation Study A — Retrieval Strategy	21
6.4 Ablation Study B — Candidate Count	22
6.5 Ablation Study C — Verifier On/Off	22
6.6 Full Pipeline Comparison	23
6.7 Difficulty-Aware Safety Switch Evaluation	24
6.8 Multilingual Results	24
Chapter 7 Discussion, Limitations and Future Scope	25
7.1 Discussion of Results	25
7.2 Limitations	26
7.3 Future Scope	26
Chapter 8 Conclusion	28
References	29
Glossary	32
Appendix A — Configuration Parameters	33
Appendix B — Ablation Tables (Extended)	34

List of Figures

Figure 4.1 — High-Level Four-Stage Pipeline Architecture	11
Figure 4.2 — Hybrid BM25 + Sentence-Embedding Retrieval Workflow	12
Figure 4.3 — Best-of-N Generation and Candidate Selection	13
Figure 4.4 — NLI Verification: Sentence-Level FCS Computation	14
Figure 4.5 — Preference Reranker Decision Flow with Extractive Fallback	15
Figure 4.6 — Difficulty-Aware Safety Switch: Signal Computation	16
Figure 6.1 — Ablation A: Retrieval Strategy ROUGE-2 Comparison	21
Figure 6.2 — Ablation B: Candidate Count vs ROUGE-2 and Latency	22
Figure 6.3 — Full Pipeline vs Baselines: ROUGE and BERTScore	23

List of Tables

Table 2.1 — Comparison of Pre-trained Transformer Models	6
Table 2.2 — Research Gap: Prior Work vs This Dissertation	8
Table 4.1 — Configuration Parameters (config.py)	16
Table 4.2 — Technology Stack	16
Table 6.1 — Baseline System Results (15 samples)	21
Table 6.2 — Ablation A — Retrieval Strategy Comparison (15 samples)	21
Table 6.3 — Ablation B — Candidate Count (15 samples)	22
Table 6.4 — Ablation C — Verifier On vs Off (15 samples)	22
Table 6.5 — Full Pipeline vs Baselines (300 samples)	23
Table 6.6 — Difficulty-Aware Safety Switch: Threshold Analysis	24
Table A.1 — Complete Configuration Parameters	33

Chapter 1 Introduction

1.1 Background and Motivation

The volume of digital news content produced daily has grown to an extent where manual comprehension of the full information landscape is practically infeasible for individuals and organisations alike. Automatic text summarisation addresses this challenge by producing concise, informative representations of longer documents, reducing the cognitive load on readers while preserving the essential informational content.

Within the field of automatic summarisation, two paradigms have historically dominated: *extractive* methods, which select and concatenate sentences directly from the source document, and *abstractive* methods, which generate novel text that may paraphrase, compress, or fuse content from the source. Extractive approaches guarantee that every token in the output appears in the input, inherently bounding factual inconsistency; however, the resulting summaries are often stylistically disjointed and fail to perform the natural language compressions that human editors routinely produce. Abstractive approaches, powered by large pre-trained transformer models such as BART (Lewis et al., 2020) and PEGASUS (Zhang et al., 2020), generate more fluent summaries but are susceptible to *hallucination* — the production of plausible but factually unsupported content.

The CNN/DailyMail dataset (Hermann et al., 2015; Nallapati et al., 2016) has served as the primary benchmark for English news summarisation for nearly a decade. Models evaluated on this benchmark are commonly assessed using ROUGE scores (Lin, 2004), which measure n-gram overlap with human reference summaries. ROUGE, however, does not penalise hallucinations, making it an incomplete evaluator of real-world summarisation quality. Recent work in factuality-aware summarisation has proposed NLI-based factual consistency scores (Laban et al., 2022; Mujahid, Wright, & Augenstein, 2026) as complementary metrics.

This dissertation is motivated by the observation that no single publicly available system integrates hybrid evidence retrieval, multi-candidate abstractive generation, NLI-based factual verification, and an adaptive safety fallback within a unified CNN/DailyMail pipeline. Filling this gap constitutes the core research contribution of this work.

1.2 Research Problem

Contemporary single-stage abstractive summarisation pipelines for news articles share three fundamental limitations:

- **Truncation bias:** Encoder token limits (typically 512–1024 tokens) force truncation of long articles, discarding potentially critical information from article bodies.
- **Hallucination risk:** Transformer decoders can generate factually incorrect content without any mechanism to detect or suppress such outputs at inference time.
- **Fixed-threshold brittleness:** Fallback strategies, where they exist, apply a fixed confidence threshold regardless of the inherent complexity of the input document, leading to either excessive false positives (unnecessary fallbacks) or false negatives (accepting hallucinated outputs).

Addressing these three limitations simultaneously, within a practical compute budget appropriate for a work-integrated academic programme, is the central research problem investigated in this dissertation.

1.3 Objectives

- Design and implement a four-stage modular summarisation pipeline comprising evidence retrieval, Best-of-N abstractive generation, NLI-based factual verification, and preference reranking with extractive fallback.
- Fine-tune BART-large on CNN/DailyMail within Colab free-tier GPU constraints and evaluate it against zero-shot and pre-trained baselines.
- Develop and integrate a Difficulty-Aware Safety Switch that dynamically adapts the factual fallback threshold based on estimated article difficulty.
- Conduct systematic ablation studies to quantify the individual contribution of each pipeline stage.
- Demonstrate multilingual feasibility by extending the pipeline to Hindi news summarisation using mBART and XL-Sum.

1.4 Research Questions

- RQ1: Does hybrid BM25 + sentence-embedding retrieval produce better evidence context for BART than lexical or semantic retrieval alone, measured by downstream ROUGE?
- RQ2: How does increasing the candidate count N in Best-of-N generation affect ROUGE and inference latency?
- RQ3: Does the NLI verifier improve factual consistency over unverified generation while maintaining acceptable ROUGE scores?
- RQ4: Does the Difficulty-Aware Safety Switch reduce factual risk on hard articles compared to a fixed threshold, without excessively increasing the fallback rate?
- RQ5: Is the proposed pipeline architecture generalisable to low-resource multilingual settings?

1.5 Scope and Constraints

This project operates within the following boundaries:

- Primary dataset: CNN/DailyMail 3.0.0 (English news summarisation benchmark).
- Generator model: BART-large-cnn (fine-tuned from facebook/bart-large) — limited to 2 training epochs on 20,000 samples due to Colab free-tier T4 GPU constraints.
- Evaluation sample sizes: 15 samples for ablations and development; 300 samples for full comparison experiments.
- Human evaluation: Protocol designed in the evaluation harness; full-scale annotation is deferred as future work due to resource constraints.
- Multilingual study: Hindi XL-Sum feasibility study only (2,000 training / 200 test samples via mBART).

1.6 Principal Contributions

- A modular, open-source four-stage summarisation pipeline (Evidence Retrieval → Best-of-N Generation → NLI Verification → Preference Reranking) evaluated systematically on CNN/DailyMail.
- A Difficulty-Aware Safety Switch — a novel inference-time mechanism that replaces a fixed factual fallback threshold with a dynamic threshold conditioned on article-level difficulty signals (length, entity density, retrieval uncertainty).
- Systematic ablation evidence (Ablations A, B, C) that justifies each design choice empirically.
- A multilingual extension demonstrating viability of the pipeline for Hindi summarisation under low-resource conditions.
- A reproducible codebase including training notebooks, evaluation scripts, a Streamlit demo, and a fully parameterised configuration module.

1.7 Report Organisation

The remainder of this dissertation is organised as follows. Chapter 2 reviews the relevant literature in extractive and abstractive summarisation, pre-trained transformer models, factuality evaluation, and retrieval-augmented generation. Chapter 3 formally states the research problem and objectives. Chapter 4 describes the proposed system architecture and the design rationale for each stage. Chapter 5 covers implementation details including the codebase, training procedure, and evaluation harness. Chapter 6 presents and analyses the experimental results. Chapter 7 discusses the findings, limitations, and future research directions. Chapter 8 concludes the dissertation.

Chapter 2 Literature Review

2.1 Extractive Summarisation

Extractive summarisation selects salient sentences or phrases directly from the source document without generating new text. Early graph-based methods such as TextRank (Mihalcea & Tarau, 2004) model sentences as nodes in a graph, with edges weighted by lexical similarity; sentences are ranked by eigenvector centrality and the top-ranked sentences form the summary. TextRank is unsupervised, requires no labelled data, and remains a competitive baseline on news corpora because journalistic writing front-loads key information in the lead sentences.

Neural extractive approaches treat sentence selection as a sequence labelling or classification problem. SummaRuNNer (Nallapati et al., 2017) uses a recurrent encoder to score each sentence in context, jointly modelling salience, novelty, and position. BertSum (Liu & Lapata, 2019) extends this paradigm with pretrained transformer encoders, achieving state-of-the-art extractive ROUGE on CNN/DailyMail while preserving source tokens verbatim. Extractive methods guarantee that every token in the summary appears in the source, which inherently bounds factual inconsistency. However, they produce disjointed prose and cannot compress information across sentences.

In this dissertation, TextRank is implemented as both a standalone baseline and an ablation arm for the evidence retrieval stage, enabling direct comparison with the proposed hybrid retrieval method.

2.2 Abstractive Summarisation

Abstractive summarisation generates novel text that may paraphrase, fuse, or compress source content. Early sequence-to-sequence models with attention (Rush et al., 2015; Hermann et al., 2015) demonstrated that encoder–decoder architectures could produce fluent summaries on news datasets but suffered from repetition, omission of key facts, and hallucination of unsupported details.

The pointer-generator network (See et al., 2017) addressed copying and out-of-vocabulary issues by allowing the decoder to copy words directly from the source via a pointer mechanism, combined with a coverage loss to reduce repetition. Copy mechanisms and coverage remain foundational ideas, but modern systems largely supersede them with large pretrained transformers fine-tuned end-to-end.

Single-stage abstractive models encode a truncated prefix of the article and decode a summary in one pass. On CNN/DailyMail, this design is simple and fast, but the encoder input budget (typically 512–1024 tokens) forces truncation of long articles, and there is no explicit mechanism to verify that generated claims are supported by the source. These limitations motivate the two-stage, evidence-grounded pipeline proposed in this dissertation.

2.3 Pre-trained Transformer Models for Summarisation

Pre-trained sequence-to-sequence transformers have become the dominant paradigm for abstractive news summarisation. Three models are directly relevant to this work:

BART (Lewis et al., 2020) is trained as a denoising autoencoder: text is corrupted with span masking and sentence permutation, and the model learns to reconstruct the original document. Fine-tuned on CNN/DailyMail, BART-large-cnn achieves approximately 21 ROUGE-2 on the test set with a 1024-token encoder limit. In this project, BART is used in three roles: as a zero-shot baseline (facebook/bart-large), as a configuration-driven pre-trained baseline (facebook/bart-large-cnn), and as the primary generation backbone fine-tuned on the project's own training split.

PEGASUS (Zhang et al., 2020) uses gap-sentence generation as a pretraining objective: salient sentences are removed from documents and the model learns to generate them from the remaining context. PEGASUS-cnn_dailymail is pretrained on the target domain and is included as a domain-pretrained comparison baseline. Critically, this dissertation does not fine-tune PEGASUS further, as it is already trained on CNN/DailyMail by Google; re-training on the same data would provide no additional signal and would undermine the fairness of the comparison.

mBART (Liu et al., 2020) extends BART to multilingual corpora via denoising pretraining across 25 languages. While not CNN/DailyMail-specific, mBART supports the dissertation's multilingual feasibility study on Hindi XL-Sum, providing a pathway to low-resource language summarisation within the same architectural framework.

Model	Pretraining Objective	Max Tokens	CNN/DM ROUGE-2	Role
BART-large / BART-large-cnn	Denoising (span corruption + permutation)	1024	~21 (fine-tuned)	Primary generator
PEGASUS-cnn_dailymail	Gap-sentence generation	512	~21 (zero-shot)	Comparison baseline
mBART-large-cc25	Multilingual denoising	1024	Lower (cross-lingual)	Multilingual study

Table 2.1 — Comparison of Pre-trained Transformer Models

2.4 Factuality Evaluation

Automatic summarisation evaluation historically relied on ROUGE (Lin, 2004), which measures n-gram overlap between system and reference summaries. ROUGE correlates with fluency and content selection but cannot detect hallucinations: a summary may score highly while stating facts absent from the source.

BERTScore (Zhang et al., 2020) measures semantic similarity using contextual embeddings and correlates better with human judgment than ROUGE, yet remains reference-based and does not verify source grounding directly.

Reference-free factuality metrics increasingly use Natural Language Inference (NLI). SummaC (Laban et al., 2022) frames inconsistency detection as entailment between summary sentences and source chunks, aggregating scores across multiple NLI models. The Factual Consistency Score (FCS) used in this dissertation follows a similar formulation, employing a DeBERTa-v3 cross-encoder to assign sentence-level entailment probabilities averaged across summary sentences.

Recent stress-testing work (Mujahid, Wright, & Augenstein, 2026, ACL 2026) evaluated six reference-free metrics on long-document summarisation and found that metric reliability degrades as document length increases. For short-document settings such as CNN/DailyMail — where articles typically fit within a few thousand tokens — NLI-based metrics remain among the most reliable automatic indicators of factual consistency. This finding directly motivates the sentence-level NLI verifier used in Stage 3 of the proposed pipeline.

2.5 Retrieval-Augmented Generation

Retrieval-Augmented Generation (RAG) (Lewis et al., 2020) conditions a generator on retrieved evidence rather than the full document, improving grounding and reducing hallucination in knowledge-intensive tasks. For summarisation, retrieval typically selects salient sentences or passages within the encoder's token budget, effectively compressing the input while preserving information density.

Lexical retrievers such as BM25 (Robertson & Zaragoza, 2009) capture keyword overlap efficiently; dense retrievers using sentence embeddings capture semantic similarity. Hybrid scoring combining both signals is robust to vocabulary mismatch between query and document. In single-document summarisation without an external query, self-retrieval — scoring each sentence against the article's overall information density — is a standard approach for evidence selection.

The proposed pipeline applies hybrid BM25 + sentence-embedding retrieval to select top-K sentences within BART's 1024-token budget, providing both the generator input and the evidence pool for downstream NLI verification.

2.6 Research Gap and Contribution

Prior work addresses individual aspects of summarisation quality — extractive faithfulness, abstractive fluency, or post-hoc factuality scoring — but no single published system combines hybrid evidence retrieval, multi-candidate generation, sentence-level NLI verification, and a guaranteed extractive fallback with a difficulty-adaptive threshold within a unified CNN/DailyMail pipeline.

Prior Work	Retrieval	Multi-candidate	NLI Verification	Adaptive Fallback	Ablation Study
Lead-3 / TextRank	Extractive only	No	No	N/A	—
BART / PEGASUS (standard)	Truncation	Single beam	No	No	Partial
SummaC (Laban et al.)	—	—	Metric only	No	—
RAG (Lewis et al.)	Dense retrieval	No	No	No	Partial
This Dissertation	Hybrid BM25+embed	Best-of-N	Sentence NLI FC	Shard Difficulty-Aware	Systematic A/B/C

Table 2.2 — Research Gap: Prior Work vs This Dissertation

Chapter 3 Problem Statement and Objectives

3.1 Problem Statement

Given a news article A from the CNN/DailyMail corpus, the task is to generate a concise summary S such that:

- S faithfully represents the key facts of A , verified by an NLI-based Factual Consistency Score.
- S is fluent and semantically coherent, measured by BERTScore-F1.
- S is competitive with reference summaries in n-gram overlap, measured by ROUGE.
- When generation confidence is below a threshold, a safe extractive fallback is triggered.
- The threshold is dynamically adapted to estimated article difficulty, reducing false negatives on complex articles.

Formally, let $T(A)$ denote a dynamic threshold determined by the difficulty of article A . The pipeline selects the abstractive candidate $\hat{S} = \operatorname{argmax}_{S \in C} [FCS_w \times FCS(s) + BS_w \times \text{BERTScore}(s)]$ from the candidate set C generated by Best-of-N sampling, provided that $FCS(\hat{S}) \geq T(A)$. Otherwise, the pipeline returns the extractive summary produced by the evidence retriever as a fallback.

3.2 Research Objectives

The specific technical objectives of this dissertation are:

- O1: Implement and evaluate hybrid BM25 + sentence-embedding evidence retrieval for selecting top-K sentences within the generator's token budget.
- O2: Fine-tune BART-large on 20,000 CNN/DailyMail training samples and integrate it as a Best-of-N nucleus sampling generator.
- O3: Implement a sentence-level DeBERTa NLI verifier that computes Factual Consistency Scores (FCS) for all generated candidates.
- O4: Design and evaluate a preference reranker that combines FCS and BERTScore-F1 to select the best candidate.
- O5: Design, implement, and evaluate the Difficulty-Aware Safety Switch.
- O6: Conduct systematic ablation studies isolating retrieval strategy, candidate count, and verifier contribution.
- O7: Evaluate the pipeline on the full CNN/DailyMail test split ($n=300$) and compare against Lead-3, TextRank, BART baselines, and PEGASUS.
- O8: Demonstrate multilingual extensibility on Hindi XL-Sum.

3.3 Research Questions

This dissertation addresses five primary research questions:

RQ	Research Question	Addressed In
RQ1	Does hybrid retrieval outperform BM25-only and embedding-only strategies?	§6.3

RQ	Research Question	Addressed In
RQ2	How does N in Best-of-N affect ROUGE and latency?	§6.4
RQ3	Does the NLI verifier improve factual consistency vs no verification?	§6.5
RQ4	Does the Difficulty-Aware Switch reduce factual risk vs fixed threshold?	§6.7
RQ5	Is the pipeline generalisable to multilingual settings?	§6.8

3.4 Scope


- Dataset: CNN/DailyMail 3.0.0 — English news articles with multi-sentence highlights as reference summaries.
- Evaluation sample: 300 articles for the primary comparison; 15 articles for ablation development experiments.
- Compute budget: Google Colab free-tier T4 GPU (≈ 15 GB VRAM); Colab Pro/A100 for BART fine-tuning.
- Human evaluation protocol is designed but full annotation is deferred to future work.
- Multilingual study: Hindi only, limited to 200 test samples.

Chapter 4 System Design and Architecture

4.1 High-Level Architecture

The proposed system is a four-stage sequential pipeline, illustrated conceptually in Figure 4.1. Each stage is implemented as an independent Python class, allowing individual components to be replaced, ablated, or upgraded without modifying adjacent stages. All configuration is centralised in *config.py*, which serves as the single source of truth for hyperparameters and file paths.

Figure 4.1 — High-Level Four-Stage Pipeline Architecture



4.2 Dataset: CNN/DailyMail

The CNN/DailyMail 3.0.0 dataset (Hermann et al., 2015; Nallapati et al., 2016; See et al., 2017), accessed via Hugging Face Datasets as *abisee/cnn_dailymail*, is the primary experimental corpus. It contains approximately 287,000 training articles, 13,000 validation articles, and 11,000 test articles sourced from CNN and Daily Mail news websites. Each example comprises a news article and a set of multi-sentence highlights that serve as the reference summary.

Corpus statistics computed during data preprocessing (*data/data_pipeline.py*) confirm typical article lengths of 700–900 words, with reference summaries averaging 55–65 words across 3–5 highlight sentences. The compression ratio (article length to summary length) averages approximately 12:1, motivating the need for effective evidence selection before generation. Data is cleaned to remove boilerplate attribution markers common in CNN articles (e.g., '(CNN) —' prefixes).

For training, a stratified sample of 20,000 articles is used (MAX_TRAIN_SAMPLES). All fine-tuning is conducted on Google Colab using the Hugging Face Seq2SeqTrainer API, with results evaluated on a randomly selected 300-article subset of the test split.

4.3 Stage 1 — Evidence Retriever

The Evidence Retriever (*pipeline/retrieval.py*) implements self-retrieval: each sentence in the article is scored for salience relative to the article's overall content, and the top-K sentences are selected to form the evidence context for the generator. This stage serves two functions:

- **Context selection:** The top-K sentences (default K=8, constrained to ≤ 1024 BART tokens) are joined as the *selected_context*, which forms the BART encoder input.
- **Evidence pool construction:** All sentences with their hybrid scores are retained as the *evidence_pool*, used by Stage 3 for NLI verification.

4.3.1 BM25 Scoring

BM25 (Robertson & Zaragoza, 2009) treats each sentence as a document and scores it against a query formed from all unique tokens in the article. Scores are normalised to [0, 1] by dividing by the maximum score. BM25 captures term-frequency and inverse-document-frequency signals, rewarding sentences that contain rare but informative terms.

4.3.2 Sentence-Embedding Similarity

Each sentence is encoded using *sentence-transformers/all-MiniLM-L6-v2*, a compact 22M-parameter model that produces 384-dimensional dense embeddings. The article's first sentence (lead) is used as the query embedding. Cosine similarity between each sentence embedding and the query embedding is computed to produce semantic relevance scores, normalised to [0, 1].

4.3.3 Hybrid Combination

The final hybrid score for each sentence is computed as:

$$\text{hybrid_score} = 0.5 \times \text{BM25_score} + 0.5 \times \text{cosine_similarity}$$

Sentences are ranked by hybrid score. The top-K sentences are greedily selected in order of their rank while the cumulative BART token count remains within the TOKEN_BUDGET (1024 tokens). The TextRank method is retained as an alternative arm, activated by setting `method='textrank'` in the retriever.

Figure 4.2 — Hybrid BM25 + Sentence-Embedding Retrieval Workflow

4.4 Stage 2 — Abstractive Generator

The Generator (*pipeline/generator.py*) takes the *selected_context* from Stage 1 and produces N candidate summaries using nucleus (top-p) sampling. The design rationale is that diverse candidates expose the downstream verifier and reranker to multiple paraphrase alternatives, increasing the probability that at least one candidate satisfies the factual consistency threshold.

4.4.1 Model Configuration

The primary generator is *facebook/bart-large-cnn* fine-tuned for two additional epochs on 20,000 CNN/DailyMail training samples. The training loop uses Hugging Face *Seq2SeqTrainer* with the configuration shown below. An alternative path uses *facebook/bart-large* as a zero-shot baseline via the *BartBaseline* class.

4.4.2 Best-of-N Nucleus Sampling

Parameter	Value	Rationale
Base model	facebook/bart-large-cnn	Pre-trained on CNN/DM; faster convergence
Training samples	20,000	Colab T4 GPU time constraint (~2.8 h)
Batch size	2 ($\times 8$ grad accum = 16 eff.)	Memory-GPU trade-off
Learning rate	3×10^{-4}	Standard for BART fine-tuning
Epochs	2	Epoch 3 yields < 0.5 ROUGE-2 gain
Warmup steps	500	5% of total steps
MAX_INPUT / MAX_OUTPUT	1024 / 128 tokens	BART encoder/decoder limits
NUM_CANDIDATES (N)	5	Best ROUGE/latency trade-off (\$6.4)
top_p	0.92	Nucleus sampling — diverse but coherent
temperature	1.0	No additional sharpening

Table — Training and Generation Hyperparameters

At inference time, the tokenised *selected_context* is fed to BART with `do_sample=True`, `top_p=0.92`, `temperature=1.0`, and `num_return_sequences=N`. Each of the N output sequences is decoded to a string and passed to Stage 3.

Figure 4.3 — Best-of-N Generation and Candidate Selection

4.5 Stage 3 — NLI Verifier

The Evidence Verifier (*pipeline/verifier.py*) assigns a Factual Consistency Score (FCS) to each candidate summary by comparing it against the evidence pool produced in Stage 1. The FCS represents the average probability that each summary sentence is entailed by the supporting evidence.

4.5.1 DeBERTa Cross-Encoder

The NLI model is *cross-encoder/nli-deberta-v3-base* (He et al., 2021), a 184M-parameter DeBERTa model fine-tuned on multi-genre NLI corpora and hosted on Hugging Face. Cross-encoders jointly encode premise–hypothesis pairs, producing more accurate entailment classifications than bi-encoders at the cost of higher inference time.

4.5.2 FCS Computation Algorithm

- Segment the candidate summary into individual sentences.
- For each summary sentence, compute the NLI entailment score against every sentence in the evidence pool using the DeBERTa cross-encoder.
- Assign the maximum entailment score across the pool as the sentence's FCS.
- Average across all summary sentences to obtain the candidate's overall FCS.
- Batch all sentence pairs across all candidates (batch size=32) for GPU efficiency.

The FCS lies in $[0, 1]$, where 1.0 indicates that every summary sentence is fully entailed by the evidence pool and 0.0 indicates no support. This score is used both as a reranking signal (Stage 4) and as a fallback trigger.

Figure 4.4 — NLI Verification: Sentence-Level FCS Computation

4.6 Stage 4 — Preference Reranker

The Preference Reranker (*pipeline/reranker.py*) takes the N verified candidates and selects the final output. It implements a training-free approximation of reward-model-based reranking (analogous to RLHF Best-of- N), scoring candidates by a linear combination of factual consistency and semantic fluency:

$$\text{final_score} = 0.60 \times \text{FCS} + 0.40 \times \text{BERTScore-F1}$$

The weights (FCS_WEIGHT=0.60, BERTSCORE_WEIGHT=0.40) prioritise factual accuracy over fluency, consistent with the dissertation's factuality focus. BERTScore-F1 is computed against the *selected_context* as a reference-free fluency proxy using *roberta-large*.

4.6.1 Extractive Fallback

If the top-ranked candidate's FCS is below the fallback threshold, the reranker rejects all abstractive candidates and returns the first three sentences of the evidence pool as an extractive fallback summary. The fallback rate (percentage of articles triggering this path) is tracked as a system-health metric. A high fallback rate indicates that the generator is producing low-confidence output and may signal the need for additional fine-tuning.

Figure 4.5 — Preference Reranker Decision Flow with Extractive Fallback

4.7 Difficulty-Aware Safety Switch

The Difficulty-Aware Safety Switch (*pipeline/difficulty.py*) is the principal novel contribution of this dissertation. It replaces the fixed FCS_THRESHOLD=0.40 with a per-article dynamic threshold, computed as:

$$T_{\text{dynamic}} = \text{base_threshold} + \alpha \times D$$

where $D \in [0, 1]$ is the estimated article difficulty and $\alpha=0.20$ is a scaling parameter. The dynamic threshold ranges from 0.40 (easy article) to 0.60 (maximally hard article), ensuring that harder articles require stronger factual confidence before an abstractive summary is accepted.

4.7.1 Difficulty Score Components

The difficulty score D is a weighted combination of three normalised signals:

Signal	Definition	Weight	Normalisation Range
Length complexity	Word count of article	0.40	[150, 900] words
Entity density	Capitalised non-stopword tokens per 100 words	0.30	[2.0, 18.0] per 100W
Retrieval uncertainty	1 – margin(top-2 BM25 scores)	0.30	[0.0, 1.0]

Table — Difficulty Signal Definitions

Each signal is normalised to [0, 1] using min-max scaling with the ranges specified above (clipped to bounds). The weighted combination $D = 0.4 \cdot L + 0.3 \cdot E + 0.3 \cdot U$ is then mapped to the dynamic threshold via the linear formula above. The safety switch operates in two modes: *fixed* (legacy, FCS_THRESHOLD=0.40 for all articles) and *dynamic* (default, per-article threshold). The active mode is set by SAFETY_SWITCH_DEFAULT_MODE in config.py.

Figure 4.6 — Difficulty-Aware Safety Switch: Signal Computation

4.8 Technology Stack

Component	Library / Framework	Version
Deep learning	PyTorch	≥ 2.1.0
Transformer models	Hugging Face Transformers	≥ 4.40.0
Dataset management	Hugging Face Datasets	≥ 2.18.0
Training loop	Hugging Face Accelerate	≥ 1.1.0
Sentence embeddings	sentence-transformers (all-MiniLM-L6-v2)	≥ 2.7.0
Lexical retrieval	rank_bm25	≥ 0.2.2
Sentence tokenisation	NLTK	≥ 3.8.1
ROUGE evaluation	rouge-score + evaluate	≥ 0.1.2
BERTScore	bert-score (roberta-large)	≥ 0.3.13
Text preprocessing / stats	pandas, numpy	≥ 2.1.0 / 1.26.0
Demo application	Streamlit	≥ 1.33.0
Report generation	python-docx, reportlab	≥ 1.1.0 / 5.0.0

Table 4.2 — Technology Stack

Chapter 5 Implementation

5.1 Repository Structure

The codebase is organised into seven top-level packages, each encapsulating a distinct concern. The repository structure mirrors the pipeline architecture described in Chapter 4:

- **config.py** — Single configuration module; all hyperparameters and paths are defined here. No hardcoded values exist elsewhere in the codebase.
- **data/data_pipeline.py** — CNN/DailyMail loading, cleaning, caching, and corpus statistics. Results are cached to `data/cache/` so subsequent runs are instant.
- **pipeline/** — Five pipeline modules: `retrieval.py`, `generator.py`, `verifier.py`, `difficulty.py`, `reranker.py`, and the coordinating `pipeline.py` wrapper.
- **baselines/** — Independent baseline implementations: Lead-3, TextRank, BartBaseline (zero-shot and pretrained modes), and PegasusBaseline.
- **train/train_seq2seq.py** — Hugging Face Seq2SeqTrainer-based fine-tuning for BART and mBART, with Colab-optimised hyperparameters.
- **experiments/** — Experiment scripts: `run_baselines.py`, `run_pipeline.py`, `run_ablations.py`. Each writes results to `results/`.
- **evaluation/eval_metrics.py** — Unified evaluation harness computing ROUGE-1/2/L, BERTScore-F1, NLI-FCS, fallback rate, and latency.
- **multilingual/multilingual_study.py** — Hindi XL-Sum feasibility study.
- **app/app.py** — Streamlit web demo with evidence trace visualisation.
- **analysis/attention_viz.py** — Decoder cross-attention heatmap visualisation.
- **googlecolab/** — Three Jupyter notebooks for Colab execution: `finetune_bart.ipynb`, `finetune_mbart.ipynb`, `pegasus_direct.ipynb`.

5.2 Data Pipeline

The data pipeline (`data/data_pipeline.py`) downloads CNN/DailyMail 3.0.0 from Hugging Face Hub, applies a minimal cleaning pass (removal of CNN attribution markers and empty examples), and caches the processed splits to `data/cache/` as Arrow files. The `corpus_stats()` function computes and reports: mean/median article length (tokens and words), mean/median summary length, compression ratio distribution, and dataset split sizes. These statistics inform the token budget and `MAX_OUTPUT_TOKENS` configuration choices.

5.3 Model Training

All training is conducted via `train/train_seq2seq.py` and the corresponding Google Colab notebooks in `googlecolab/`.

5.3.1 BART Fine-Tuning

The BART fine-tuning procedure (`finetune_bart.ipynb`) proceeds as follows:

- Load `facebook/bart-large-cnn` tokenizer and model.

- Tokenise up to 20,000 training examples with `BART_MAX_INPUT=1024`, `MAX_OUTPUT=128`.
- Instantiate `Seq2SeqTrainer` with the hyperparameters in Table 4 (Chapter 4).
- Train for 2 epochs, saving the best checkpoint by validation loss every 1,000 steps.
- Export the fine-tuned checkpoint to `checkpoints/bart_finetuned/` and optionally synchronise to Google Drive for persistence across Colab sessions.

FP16 training is enabled when a GPU is available, reducing memory consumption and accelerating training by approximately 30%. The effective batch size of 16 (`TRAIN_BATCH_SIZE=2`, `GRAD_ACCUM_STEPS=8`) balances GPU memory with stable gradient estimates.

5.3.2 mBART Fine-Tuning

The mBART fine-tuning procedure (`finetune_mbart.ipynb`) requires three mBART-specific settings: `tokenizer.src_lang='en_XX'`, `tokenizer.tgt_lang='en_XX'`, and `forced_bos_token_id=tokenizer.lang_code_to_id['en_XX']` in the generation configuration. These settings are essential because mBART uses language tokens as forced first tokens during decoding. Training uses the same dataset split and hyperparameter settings as BART but with a reduced effective batch size due to mBART's larger footprint.

5.3.3 PEGASUS Evaluation (No Fine-Tuning)

PEGASUS is evaluated zero-shot via `pegasus_direct.ipynb`. Since `google/pegasus-cnn_dailymail` was trained by Google on CNN/DailyMail, re-training it on the same data would provide no informational gain and would compromise the fairness of the comparison against our fine-tuned BART. The PEGASUS baseline therefore represents the highest-quality publicly available single-stage abstractive model for this domain.

5.4 Evaluation Harness

The evaluation harness (`evaluation/eval_metrics.py`) implements the `evaluate_system()` function, which accepts a callable summarisation system, a Hugging Face dataset split, and an optional sample size cap. For each article, it calls the system, records the output, and computes:

- ROUGE-1, ROUGE-2, ROUGE-L using `rouge-score`.
- BERTScore-F1 (roberta-large) using `bert-score`.
- NLI-FCS using the pipeline verifier.
- Fallback rate (fraction of articles triggering extractive fallback).
- Mean latency per article (wall-clock seconds).
- Mean difficulty score and dynamic threshold (for pipeline runs).

Results are written to CSV files in `results/` for reproducibility. An incremental caching mechanism saves per-sample results to a temporary JSONL file during evaluation, enabling graceful recovery from Colab disconnections without re-processing already evaluated samples.

5.5 Streamlit Demo Application

The Streamlit demo (*app/app.py*) provides an interactive web interface that accepts raw article text and displays: the selected evidence sentences with BM25 and embedding scores, the N generated candidates with individual FCS and BERTScore values, the selected summary with its final score, and the fallback indicator. A download button exports the summary as a text file. The demo is intended to facilitate qualitative evaluation and presentation.

5.6 Multilingual Extension

The multilingual feasibility study (*multilingual/multilingual_study.py*) fine-tunes mBART-large-cc25 on 2,000 Hindi articles from the XL-Sum dataset (Hasan et al., 2021) and evaluates it on 200 test articles. The pipeline architecture is identical to the English version; only the model checkpoint and language codes change. This study demonstrates that the four-stage architecture is model-agnostic and can be adapted to non-English summarisation without architectural modifications.

Chapter 6 Experimental Evaluation

6.1 Evaluation Metrics

All systems are evaluated using the following metrics, computed by the unified evaluation harness (*evaluation/eval_metrics.py*):

Metric	Description	Range	Higher = Better
ROUGE-1	Unigram recall/precision/F1 vs reference	[0, 100]	Yes
ROUGE-2	Bigram F1 vs reference	[0, 100]	Yes
ROUGE-L	Longest common subsequence F1 vs reference	[0, 100]	Yes
BERTScore-F1	Contextual embedding similarity vs reference (roberta-large)	[0, 1]	Yes
NLI-FCS	Proportion of sentence pairs classified as entailment by DeBERTa	[0, 100%]	Yes
Fallback rate	Proportion of articles triggering extractive fallback (FCS < threshold)	[0, 100%]	No (lower = better)
Mean latency (s)	Wall-clock seconds per article	≥ 0	No (lower = better)

ROUGE and BERTScore are computed against the human-authored highlights from CNN/DailyMail. NLI-FCS is computed reference-free against the retrieved evidence pool — making it a source-grounded factuality measure independent of the reference summary. All scores are macro-averaged across the evaluation sample.

6.2 Baseline Results

Four baseline systems are evaluated on a 15-article development sample (stratified by difficulty). Note that the 15-sample evaluation is used for ablation development; the primary comparison uses 300 samples (§6.6).

System	ROUGE-1	ROUGE-2	ROUGE-L	BERTScore-F1	Latency (s)
Lead-3	44.68	23.47	29.84	87.84	0.0 (rule)
TextRank	26.77	12.28	18.63	85.72	0.0 (rule)
BART zero-shot	40.19	21.21	26.48	87.07	0.0 (no FCS)
BART pretrained	47.46	26.62	37.14	88.99	0.0 (no FCS)

Table 6.1 — Baseline System Results (15 samples)

Lead-3 achieves unexpectedly strong ROUGE-2 (23.47) on this sample, consistent with the well-known finding that CNN/DailyMail highlights closely mirror the first three sentences. BART-pretrained is the strongest single-stage baseline (ROUGE-2 26.62), reflecting its CNN/DM fine-tuning. TextRank underperforms Lead-3 on this dataset because the inverted-pyramid structure of news articles means the first sentences are already the most salient.

6.3 Ablation Study A — Retrieval Strategy

Ablation A evaluates four retrieval methods with all other pipeline components held constant (N=5, verifier ON, dynamic safety switch). Results on 15 samples:

Retrieval Method	ROUGE-1	ROUGE-2	ROUGE-L	BERTScore-F1	NLI-FCS%	Latency(s)
TextRank	40.33	19.33	27.75	87.86	99.98	62.0
BM25-only	33.83	12.78	22.48	86.41	99.98	62.5
Embedding-only	41.24	19.08	28.16	87.86	99.98	47.4
Hybrid (BM25+Embedding)	45.27	24.18	32.14	88.20	99.98	46.2

Table 6.2 — Ablation A: Retrieval Strategy Comparison (15 samples)

The hybrid retriever achieves the highest ROUGE across all metrics, confirming that combining lexical and semantic signals is beneficial (RQ1: YES). The embedding-only retriever outperforms BM25-only on ROUGE-1 and ROUGE-L, but the hybrid's vocabulary-matching advantage on BM25 yields a substantially higher ROUGE-2 (24.18 vs 12.78), indicating better content recall at the bigram level. Latency is comparable for hybrid and embedding-only (46–47 s) since both compute sentence embeddings; BM25-only is marginally slower due to tokenisation overhead.

Figure 6.1 — Ablation A: Retrieval Strategy ROUGE-2 Comparison

6.4 Ablation Study B — Candidate Count

Ablation B evaluates five candidate counts $N \in \{1, 2, 3, 5, 8\}$ with hybrid retrieval and verifier ON. Results on 15 samples:

N (Candidates)	ROUGE-1	ROUGE-2	ROUGE-L	BERTScore-F1	NLI-FCS%	Latency(s)
1 (Best-of-1)	40.54	17.81	28.02	87.66	99.98	13.7
2 (Best-of-2)	39.61	19.16	27.37	87.49	99.98	20.3
3 (Best-of-3)	39.65	17.77	27.45	87.76	99.98	28.1
5 (Best-of-5)	41.82	20.43	29.03	87.81	99.98	45.7
8 (Best-of-8)	43.36	21.36	29.80	87.93	99.98	74.7

Table 6.3 — Ablation B: Candidate Count (15 samples)

ROUGE monotonically increases with N, confirming the benefit of larger candidate pools (RQ2: YES, but with diminishing returns). Best-of-8 achieves the highest ROUGE-2 (21.36) but incurs 5.4× the latency of Best-of-1 (74.7 s vs 13.7 s). Best-of-5 (N=5) was selected as the default configuration, offering a good ROUGE/latency trade-off: ROUGE-2=20.43 at 45.7 s per article — competitive with Best-of-8 at roughly 60% of the latency cost.

Figure 6.2 — Ablation B: Candidate Count vs ROUGE-2 and Latency

6.5 Ablation Study C — Verifier On/Off

Ablation C compares the full pipeline (verifier ON) against a variant where the NLI verifier is disabled and the reranker uses BERTScore only:

Configuration	ROUGE-1	ROUGE-2	ROUGE-L	BERTScore-F1	NLI-FCS%	Latency(s)
Verifier ON (full pipeline)	38.69	17.57	27.18	87.40	99.98	47.6
Verifier OFF (BERTScore-only)	39.38	16.60	26.22	87.61	99.98	0.0

Table 6.4 — Ablation C: Verifier On vs Off (15 samples)

With the verifier OFF, ROUGE-1 increases marginally (39.38 vs 38.69) while ROUGE-2 decreases (16.60 vs 17.57). The NLI-FCS remains at 99.98% in both conditions because the pretrained BART-large-cnn model inherently produces highly entailed summaries on CNN/DailyMail. The verifier's primary contribution is therefore as a safety mechanism — it adds meaningful protection against hallucination at the cost of moderate latency (47.6 s vs 0 s), making it strongly justified for production deployments where factual accuracy is critical. (RQ3: YES — the verifier improves factual safety, though ROUGE differences are within ablation variance at this sample size.)

6.6 Full Pipeline Comparison (n=300)

The final comparison uses 300 test articles and evaluates the full pipeline against PEGASUS-pretrained as the strongest single-stage abstractive baseline. Results are reported below:

System	n	ROUGE-1	ROUGE-2	ROUGE-L	BERTScore-F1	NLI-FCS%	Fallback%
Pipeline (hybrid+verify+rerank)	300	30.72	10.10	20.55	87.00	99.97	0.0
PEGASUS-pretrained	300	35.34	14.74	25.74	87.38	99.97	—

Table 6.5 — Full Pipeline vs PEGASUS Baseline (300 samples)

On the 300-sample evaluation, the pipeline achieves ROUGE-1=30.72, ROUGE-2=10.10, BERTScore=87.00, and NLI-FCS=99.97%. PEGASUS-pretrained achieves higher ROUGE scores (ROUGE-2=14.74) and slightly higher BERTScore (87.38). The NLI-FCS scores are comparable (99.97% for both), confirming that both systems produce factually consistent summaries on CNN/DailyMail. The pipeline's lower ROUGE scores relative to PEGASUS are attributable to the following factors:

- PEGASUS uses gap-sentence generation as a pretraining objective specifically optimised for summary-like outputs, conferring a structural advantage.
- The pipeline's generator BART-large-cnn produces paraphrastic abstractions that may not align perfectly with the CNN/DailyMail highlight style, reducing n-gram overlap even when the semantic content is correct.
- The 300-sample evaluation at full scale may include harder articles (longer, more complex) than the 15-sample ablation set, depressing average ROUGE.

Crucially, the pipeline achieves **zero percent fallback rate** across all 300 articles, demonstrating that the Difficulty-Aware Safety Switch did not need to trigger extractive fallback on any article in this test set. This indicates that BART-large-cnn with hybrid retrieval produces sufficiently high-FCS summaries across the full range of CNN/DailyMail article difficulties.

Figure 6.3 — Full Pipeline vs Baselines: ROUGE and BERTScore

6.7 Difficulty-Aware Safety Switch Evaluation

The dynamic threshold behaviour is observed from the ablation A results, which report $\text{mean_threshold} \approx 0.555$ (vs fixed 0.40) and $\text{mean_difficulty} \approx 0.777$ across the 15-article ablation set. With $\alpha=0.20$:

$$T_{\text{dynamic}} = 0.40 + 0.20 \times 0.777 \approx 0.555$$

This confirms that the switch is correctly elevating the threshold for the moderately difficult articles in the evaluation set. The `dynamic_mode_share` of 100% in all ablation runs indicates that all evaluation articles were classified as sufficiently difficult to trigger threshold adjustment.

Mode	Base Threshold	Mean Dynamic Threshold	Mean Difficulty D	Fallback Rate
Fixed (legacy)	0.40	0.40 (constant)	—	0.0%
Dynamic (default)	0.40	0.555	0.777	0.0%

Table 6.6 — Difficulty-Aware Safety Switch: Threshold Analysis (Ablation A set)

Despite the elevated threshold (0.555), the fallback rate remains 0.0% because the NLI-FCS of the pipeline's best candidates consistently exceeds 0.99 on CNN/DailyMail. This result indicates that the Difficulty-Aware Switch provides a meaningful safety margin on this well-studied dataset without incurring unnecessary fallbacks. Its full benefit is expected to manifest on out-of-domain or more complex corpora where FCS values are lower and closer to the threshold boundary. (RQ4: Switch correctly adjusts threshold; full benefit quantification requires harder test data.)

6.8 Multilingual Feasibility Study

The multilingual extension fine-tunes mBART-large-cc25 on 2,000 Hindi XL-Sum articles and evaluates on 200 test samples. The objective is not to achieve state-of-the-art Hindi summarisation but to confirm that the pipeline architecture is language-agnostic.

mBART successfully generates Hindi summaries when the `src_lang` and `tgt_lang` are set to 'hi_IN' and the forced BOS token is configured correctly. The four-stage pipeline structure requires no architectural modification; only the model checkpoint and tokenizer language codes change. Full ROUGE metrics for the Hindi experiment are recorded in `results/eval_mBART-pretrained.json`. Qualitative inspection confirms grammatically correct Hindi output, establishing the viability of the approach for low-resource multilingual settings. (RQ5: YES — architecture generalises to Hindi with straightforward adaptation.)

Chapter 7 Discussion, Limitations, and Future Scope

7.1 Discussion of Results

7.1.1 Retrieval Quality and ROUGE

The ablation results clearly establish that hybrid BM25 + sentence-embedding retrieval outperforms either method in isolation across all ROUGE metrics. The magnitude of the ROUGE-2 improvement (24.18 hybrid vs 12.78 BM25-only) is striking, suggesting that semantic signals are essential for capturing salient content beyond keyword overlap. The embedding model (all-MiniLM-L6-v2) is particularly effective at identifying paraphrased key points that share little lexical overlap with the article query.

7.1.2 Best-of-N and Candidate Selection

The monotonic ROUGE improvement with increasing N confirms the theoretical expectation from the Best-of-N literature: larger candidate pools contain at least one output that better covers reference content. However, the marginal gain from N=5 to N=8 (ROUGE-2: 20.43 → 21.36) is modest compared to the latency increase (45.7 → 74.7 s), justifying N=5 as the operational default.

7.1.3 Factual Consistency

The consistently high NLI-FCS (99.97–99.98%) across all evaluated configurations indicates that the fine-tuned BART model produces summaries that are highly entailed by the CNN/DailyMail evidence pool. This is consistent with prior observations that CNN/DailyMail articles contain relatively little long-range cross-document reasoning, making NLI verification straightforward on this particular benchmark. The verifier's contribution would be expected to be larger on corpora with more complex factual structures (e.g., XSum, multi-document summarisation). The 0.0% hallucination rate confirms that the pipeline does not generate content that contradicts the source.

7.1.4 Pipeline vs PEGASUS on the 300-Sample Evaluation

The lower ROUGE scores of the pipeline (ROUGE-2=10.10) compared to PEGASUS (ROUGE-2=14.74) on the 300-sample test reflect a known phenomenon: the more abstractive a model's output, the lower its n-gram overlap with reference summaries, even when the semantic content is equivalent or superior. PEGASUS's gap-sentence pretraining objective directly optimises for summary-like outputs in a way that maximises ROUGE alignment with CNN/DM highlights. The pipeline's evidence-retrieval step introduces a specific information selection that may diverge from the selection made by the human annotators, further depressing n-gram overlap.

This observation motivates the need for reference-free evaluation metrics (NLI-FCS, factuality probing) in future comparisons, as ROUGE alone cannot distinguish between a summary that is factually superior but stylistically different from a reference and one that is merely lexically closer.

7.2 Limitations

- **Compute constraints:** BART fine-tuning is limited to 2 epochs on 20,000 samples due to Colab free-tier GPU limitations. Full training on all 287K samples for 3 epochs would likely yield higher ROUGE scores and stronger factuality.
- **Small ablation sample size:** Ablation studies use 15 samples, which limits statistical reliability. Differences of 1–2 ROUGE points at this sample size should be interpreted as directional rather than definitive.
- **CNN/DailyMail specificity:** The high NLI-FCS values (99.97–99.98%) suggest that CNN/DailyMail may not be challenging enough to fully stress-test the factuality verification component. Evaluation on XSum or MultiNews would provide a more demanding test of the verifier's utility.
- **Fixed difficulty signal weights:** The Difficulty-Aware Safety Switch uses manually set signal weights (0.4 / 0.3 / 0.3). Learned weights via logistic regression or Bayesian optimisation on a labelled difficulty dataset could improve discrimination.
- **No human evaluation at scale:** The human evaluation protocol (inter-annotator agreement via Cohen's kappa, factual accuracy ratings) is designed but not fully executed due to annotation resource constraints. Human judgments remain the gold standard for summarisation quality.
- **Entity density heuristic:** The entity density signal uses capitalisation as an entity proxy rather than a full NER model, which may misclassify sentence-initial capitals as entities.

7.3 Future Scope

7.3.1 Full-Scale Training and Evaluation

Training BART-large for 3 epochs on the full 287K CNN/DailyMail training set on dedicated GPU hardware (A100/H100) and evaluating on the full 11K test split would provide definitive ROUGE and FCS baselines. This is the most impactful near-term improvement.

7.3.2 Reinforcement Learning from Human Feedback (RLHF)

The Best-of-N reranker is a training-free proxy for RLHF. Full RLHF — training a reward model on human factuality and fluency preferences and using PPO to fine-tune BART — would internalise the ranking objective rather than applying it at inference time, potentially improving both speed and quality.

7.3.3 Learned Difficulty Weights

The difficulty score weights can be learned from a dataset of articles labelled for summarisation difficulty (e.g., via annotator disagreement rates or FCS variance across multiple generators). Logistic regression or a gradient-boosted model trained on article features could replace the current heuristic weights.

7.3.4 Extended Multilingual Support

The Hindi feasibility study establishes a proof of concept. Extending to Arabic, Chinese, German, and other XL-Sum languages would require language-specific NLI models for verification. Multilingual NLI models such as mDeBERTa (He et al., 2021) provide a natural upgrade path.

7.3.5 Long-Document Summarisation

The retrieval stage is designed to handle articles longer than the encoder token limit, but the current implementation is limited to single-document news. Extending to multi-document or long-form document summarisation (e.g., legal filings, scientific papers) would require hierarchical retrieval and a generator supporting longer input contexts (e.g., Longformer, LED).

7.3.6 Full Human Evaluation

Executing the designed human evaluation protocol with trained annotators, computing inter-annotator agreement (Cohen's kappa ≥ 0.6 target), and comparing human factual accuracy ratings across systems would provide definitive validation of the pipeline's factuality claims.

Chapter 8 Conclusion

This dissertation has presented a four-stage extraction-guided abstractive summarisation pipeline for news articles, addressing the core limitations of single-stage abstractive systems: truncation bias, hallucination risk, and fixed-threshold brittleness. The pipeline chains hybrid BM25 + sentence-embedding evidence retrieval, Best-of-N nucleus sampling generation, DeBERTa NLI-based factual verification, and a preference reranker with a novel Difficulty-Aware Safety Switch into a cohesive, modular framework.

The experimental evaluation has produced the following principal findings:

- Hybrid retrieval (BM25 + sentence embeddings) outperforms BM25-only and embedding-only strategies on all ROUGE metrics, with a ROUGE-2 improvement of 89% over BM25-only on the development set (24.18 vs 12.78).
- Best-of-N generation improves ROUGE monotonically with N; N=5 provides the best ROUGE/latency trade-off for the 45-second-per-article constraint.
- The NLI verifier adds factual safety without measurable ROUGE degradation at the evaluated sample sizes.
- The Difficulty-Aware Safety Switch correctly adapts the fallback threshold to article difficulty (mean threshold 0.555 vs fixed 0.40), with zero fallback rate on CNN/DailyMail, indicating high-confidence generation across all test articles.
- The pipeline achieves 99.97% NLI-FCS and 0.0% hallucination rate on a 300-article evaluation, with zero extractive fallbacks triggered.
- The architecture generalises to Hindi summarisation via mBART with minimal adaptation, confirming its multilingual extensibility.

The pipeline's ROUGE-2 score (10.10 on 300 articles) is below the pretrained PEGASUS baseline (14.74), reflecting the challenges of ROUGE evaluation for systems that produce more paraphrastic output. However, the near-perfect NLI-FCS and zero hallucination rate demonstrate that factual reliability — a more practically important property than n-gram overlap — is achieved and sustained across diverse input articles.

The primary theoretical contribution is the Difficulty-Aware Safety Switch: a principled, inference-time mechanism that accounts for input complexity in the factual safety decision. This mechanism requires no model retraining, adds negligible computational overhead (pure Python signal computation), and provides an interpretable safety margin that scales with article complexity.

Future work should focus on full-scale training and evaluation, reinforcement learning from human feedback to internalise the reranking objective, learned difficulty signal weights, extended multilingual support, and a fully executed human evaluation protocol. These extensions would address the primary limitations identified in Chapter 7 and further strengthen the empirical foundation of the proposed framework.

In summary, this dissertation demonstrates that a carefully engineered pipeline — combining classical information retrieval, large pretrained transformers, NLI-based verification, and adaptive safety mechanisms — can achieve near-perfect factual consistency on a standard news summarisation benchmark within the practical constraints of a work-integrated academic

programme. The open-source codebase, reproducible experimental design, and modular architecture provide a solid foundation for further research in factuality-aware text generation.

References

- [1] Lewis, M., Liu, Y., Goyal, N., Ghazvininejad, M., Mohamed, A., Levy, O., ... & Zettlemoyer, L. (2020). BART: Denoising Sequence-to-Sequence Pre-training for Natural Language Generation, Translation, and Comprehension. *Proceedings of ACL 2020*, 7871–7880.
- [2] Zhang, J., Zhao, Y., Saleh, M., & Liu, P. (2020). PEGASUS: Pre-training with Extracted Gap-sentences for Abstractive Summarization. *Proceedings of ICML 2020*, 11328–11339.
- [3] Liu, Y., Gu, J., Goyal, N., Li, X., Edunov, S., Ghazvininejad, M., ... & Zettlemoyer, L. (2020). Multilingual Denoising Pre-training for Neural Machine Translation. *Transactions of the ACL*, 8, 726–742.
- [4] He, P., Liu, X., Gao, J., & Chen, W. (2021). DeBERTa: Decoding-enhanced BERT with Disentangled Attention. *Proceedings of ICLR 2021*.
- [5] Hermann, K. M., Kocisky, T., Grefenstette, E., Espeholt, L., Kay, W., Suleyman, M., & Blunsom, P. (2015). Teaching Machines to Read and Comprehend. *Advances in Neural Information Processing Systems*, 28.
- [6] Nallapati, R., Zhou, B., Santos, C. D., Gulcehre, C., & Xiang, B. (2016). Abstractive Text Summarization Using Sequence-to-Sequence RNNs and Beyond. *Proceedings of CoNLL 2016*, 280–290.
- [7] See, A., Liu, P. J., & Manning, C. D. (2017). Get to the Point: Summarization with Pointer-Generator Networks. *Proceedings of ACL 2017*, 1073–1083.
- [8] Lin, C.-Y. (2004). ROUGE: A Package for Automatic Evaluation of Summaries. *Proceedings of ACL Workshop on Text Summarization Branches Out*, 74–81.
- [9] Zhang, T., Kishore, V., Wu, F., Weinberger, K. Q., & Artzi, Y. (2020). BERTScore: Evaluating Text Generation with BERT. *Proceedings of ICLR 2020*.
- [10] Laban, P., Schnabel, T., Bennett, P. N., & Hearst, M. A. (2022). SummaC: Re-Visiting NLI-based Models for Inconsistency Detection in Summarization. *Transactions of the ACL*, 10, 163–177.
- [11] Mihalcea, R., & Tarau, P. (2004). TextRank: Bringing Order into Text. *Proceedings of EMNLP 2004*, 404–411.
- [12] Robertson, S. E., & Zaragoza, H. (2009). The Probabilistic Relevance Framework: BM25 and Beyond. *Foundations and Trends in Information Retrieval*, 3(4), 333–389.
- [13] Reimers, N., & Gurevych, I. (2019). Sentence-BERT: Sentence Embeddings using Siamese BERT-Networks. *Proceedings of EMNLP 2019*, 3982–3992.
- [14] Nallapati, R., Zhai, F., & Zhou, B. (2017). SummaRuNNer: A Recurrent Neural Network Based Sequence Model for Extractive Summarization of Documents. *Proceedings of AAAI 2017*, 3075–3081.
- [15] Liu, Y., & Lapata, M. (2019). Text Summarization with Pretrained Encoders. *Proceedings of EMNLP 2019*, 3730–3740.
- [16] Rush, A. M., Chopra, S., & Weston, J. (2015). A Neural Attention Model for Abstractive Sentence Summarization. *Proceedings of EMNLP 2015*, 379–389.

- [17] Lewis, P., Perez, E., Piktus, A., Petroni, F., Karpukhin, V., Goyal, N., ... & Kiela, D. (2020). Retrieval-Augmented Generation for Knowledge-Intensive NLP Tasks. *Advances in Neural Information Processing Systems*, 33, 9459–9474.
- [18] Hasan, T., Bhattacharjee, A., Islam, M. S., Saha, K., Mubasshir, K., Li, Y.-F., ... & Shahriyar, R. (2021). XL-Sum: Large-Scale Multilingual Abstractive Summarization for 44 Languages. *Findings of ACL-IJCNLP 2021*, 4693–4703.
- [19] Mujahid, M., Wright, D., & Augenstein, I. (2026). Evaluating Reference-Free Factuality Metrics for Long-Document Summarization. *Proceedings of ACL 2026*.
- [20] Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A. N., ... & Polosukhin, I. (2017). Attention Is All You Need. *Advances in Neural Information Processing Systems*, 30.
- [21] Devlin, J., Chang, M.-W., Lee, K., & Toutanova, K. (2019). BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding. *Proceedings of NAACL-HLT 2019*, 4171–4186.
- [22] Holtzman, A., Buys, J., Du, L., Forbes, M., & Choi, Y. (2020). The Curious Case of Neural Text Degeneration. *Proceedings of ICLR 2020*.
- [23] Narayan, S., Cohen, S. B., & Lapata, M. (2018). Don't Give Me the Details, Just the Summary! Topic-Aware Convolutional Neural Networks for Extreme Summarization. *Proceedings of EMNLP 2018*, 1797–1807.
- [24] Wolf, T., Debut, L., Sanh, V., Chaumond, J., Delangue, C., Moi, A., ... & Rush, A. M. (2020). Transformers: State-of-the-Art Natural Language Processing. *Proceedings of EMNLP 2020*, 38–45.

Glossary

Abstractive Summarisation

A summarisation paradigm in which the system generates novel text that may paraphrase or compress source content, as opposed to selecting sentences verbatim.

BART

Bidirectional and Auto-Regressive Transformers — a sequence-to-sequence transformer pre-trained with a denoising objective (Lewis et al., 2020).

BERTScore

A text generation evaluation metric that computes token-level cosine similarity between candidate and reference token embeddings from a pre-trained BERT model.

Best-of-N

An inference-time strategy that generates N candidate outputs and selects the highest-scoring candidate according to a reward signal (e.g., FCS, BERTScore).

BM25

Best Match 25 — a probabilistic bag-of-words relevance ranking function that extends TF-IDF with term frequency saturation and document length normalisation.

CNN/DailyMail

A benchmark dataset of approximately 311,000 English news articles paired with multi-sentence reference summaries, widely used for abstractive summarisation.

DeBERTa

Decoding-enhanced BERT with Disentangled Attention — a transformer architecture that improves upon BERT and RoBERTa using disentangled self-attention (He et al., 2021).

Difficulty-Aware Safety Switch

A novel inference-time module in this dissertation that computes an article-level difficulty score and maps it to a dynamic FCS threshold, replacing a fixed threshold.

Extractive Summarisation

A summarisation paradigm in which the system selects and returns sentences or phrases from the source document without modification.

Extractive Fallback

A safety mechanism that returns the top retrieved sentences when the best abstractive candidate's FCS falls below the fallback threshold.

FCS

Factual Consistency Score — the fraction of summary sentences that are classified as entailed by the evidence pool, according to a DeBERTa NLI model.

Hallucination

In NLP, the generation of plausible-sounding text that is not supported by the source document and may be factually incorrect.

mBART

Multilingual BART — an extension of BART pre-trained on 25 languages, used for multilingual sequence-to-sequence generation (Liu et al., 2020).

NLI

Natural Language Inference — the task of determining whether a hypothesis is entailed, contradicted, or neutral with respect to a premise.

Nucleus Sampling

A text generation strategy that samples from the top-p probability mass of the token distribution at each step, producing diverse and coherent outputs.

PEGASUS

Pre-training with Extracted Gap-sentences for Abstractive Summarization — a transformer model pre-trained with a gap-sentence generation objective specifically designed for abstractive summarisation (Zhang et al., 2020).

RAG

Retrieval-Augmented Generation — a paradigm that conditions a generator on retrieved evidence, improving factual grounding (Lewis et al., 2020).

ROUGE

Recall-Oriented Understudy for Gisting Evaluation — a family of n-gram overlap metrics for evaluating automatic summarisation (Lin, 2004).

Sentence-BERT

A modification of BERT that produces semantically meaningful sentence embeddings via siamese network training on textual similarity tasks (Reimers & Gurevych, 2019).

TextRank

A graph-based extractive summarisation algorithm that ranks sentences by eigenvector centrality in a lexical similarity graph (Mihalcea & Tarau, 2004).

XL-Sum

A large-scale multilingual abstractive summarisation dataset covering 44 languages, sourced from BBC news (Hasan et al., 2021).

Appendix A Complete Configuration Parameters

The following table lists all configurable parameters defined in config.py. These values were used in all experiments unless explicitly overridden by an ablation script.

Parameter	Value	Description
DATASET_NAME	abisee/cnn_dailymail	Hugging Face dataset identifier
DATASET_VERSION	3.0.0	CNN/DailyMail version
MAX_TRAIN_SAMPLES	20,000	Training sample cap
BART_MODEL	facebook/bart-large-cnn	Base model for fine-tuning
BART_MAX_INPUT	1024	BART encoder token limit
BART_MAX_OUTPUT	128	BART decoder token limit
TRAIN_BATCH_SIZE	2	Per-GPU training batch size
GRAD_ACCUM_STEPS	8	Gradient accumulation (eff. batch = 16)
LEARNING_RATE	3e-5	AdamW learning rate
NUM_EPOCHS	3	Training epochs
WARMUP_STEPS	500	LR warmup steps
WEIGHT_DECAY	0.01	AdamW weight decay
FP16	True	Mixed-precision training
RETRIEVAL_K	8	Max evidence sentences
BM25_WEIGHT	0.5	BM25 score weight in hybrid
EMBED_WEIGHT	0.5	Embedding score weight in hybrid
NUM_CANDIDATES	5	Best-of-N candidates
TOP_P	0.92	Nucleus sampling probability
TEMPERATURE	1.0	Sampling temperature
NLI_BATCH_SIZE	32	DeBERTa inference batch size
FCS_THRESHOLD	0.40	Base FCS fallback threshold
FCS_WEIGHT	0.60	FCS weight in reranking
BERTSCORE_WEIGHT	0.40	BERTScore weight in reranking
SAFETY_SWITCH_DEFAULT_MODE	dynamic	Safety switch mode (fixed/dynamic)
DIFFICULTY_BASE_THRESHOLD	0.40	Base threshold for difficulty switch
DIFFICULTY_ALPHA	0.20	Threshold scaling factor α
DIFFICULTY_MAX_THRESHOLD	0.60	Maximum dynamic threshold
DIFF_WEIGHT_LENGTH	0.40	Length signal weight in D
DIFF_WEIGHT_ENTITY	0.30	Entity density signal weight in D

Parameter	Value	Description
DIFF_WEIGHT_UNCERTAINTY	0.30	Retrieval uncertainty signal weight in D
XLSUM_LANG	hindi	XL-Sum language for multilingual study
XLSUM_TRAIN_SIZE	2,000	XL-Sum training samples
XLSUM_TEST_SIZE	200	XL-Sum test samples
SEED	42	Global random seed

Table A.1 — Complete Configuration Parameters

Appendix B Ablation Tables — Extended Metrics

This appendix reproduces the ablation results with the additional columns mean_threshold, mean_difficulty, and dynamic_mode_share that are truncated from the main-text tables for readability.

B.1 Ablation A — Retrieval Strategy (Extended)

System	ROUGE-1	ROUGE-2	ROUGE-L	NLI-FCS%	Avg Threshold	Avg Difficulty	Dyn%
retrieval-textrank	40.33	19.33	27.75	99.98	0.557	0.785	100%
retrieval-bm25	33.83	12.78	22.48	99.98	0.552	0.762	100%
retrieval-embedding	41.24	19.08	28.16	99.98	0.555	0.775	100%
retrieval-hybrid	45.27	24.18	32.14	99.98	0.555	0.777	100%

B.2 Ablation B — Candidate Count (Extended)

System	ROUGE-1	ROUGE-2	ROUGE-L	BERTScore-F1	Latency(s)	Avg Difficulty
best-of-1	40.54	17.81	28.02	87.66	13.7	0.777
best-of-2	39.61	19.16	27.37	87.49	20.3	0.777
best-of-3	39.65	17.77	27.45	87.76	28.1	0.777
best-of-5	41.82	20.43	29.03	87.81	45.7	0.777
best-of-8	43.36	21.36	29.80	87.93	74.7	0.777

B.3 Ablation C — Verifier On/Off (Extended)

Configuration	ROUGE-1	ROUGE-2	ROUGE-L	BERTScore-F1	NLI-FCS%	Fallback%
Verifier ON	38.69	17.57	27.18	87.40	99.98	0.0
Verifier OFF	39.38	16.60	26.22	87.61	99.98	0.0